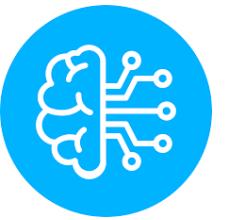


Deep Learning

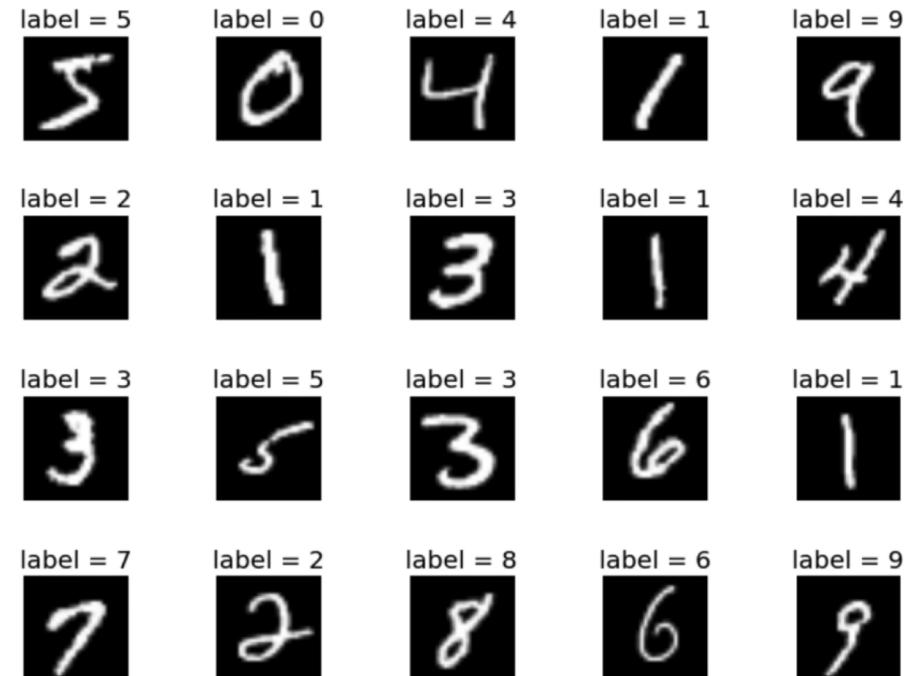
ABNANE Ibtissam



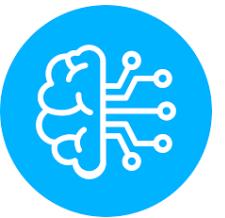
Transfer Learning

Good deep learning models are dependednt on:

- Model architecture
 - Number/size of filters
 - Number of convolution layers
- Quality of training
 - Time
 - Good data
 - Computational power



Transfer learning: Taking the knowledge resulted from the completion of one task, and implement it to achieve a new task



Transfer Learning

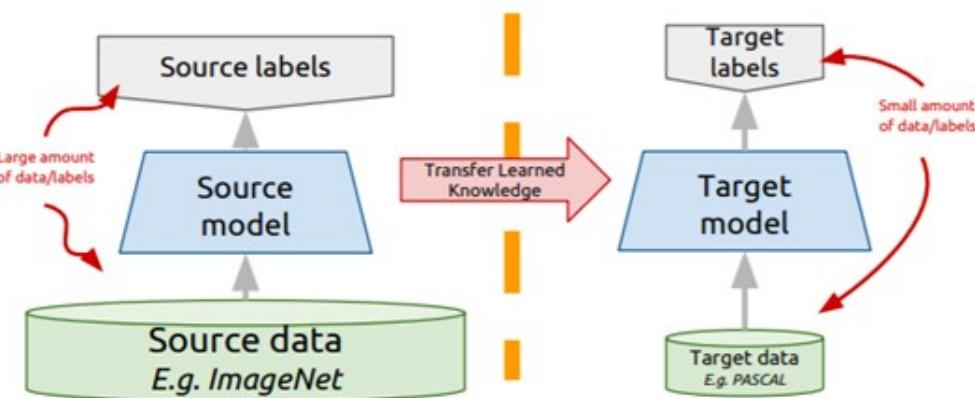
Transfer learning: idea

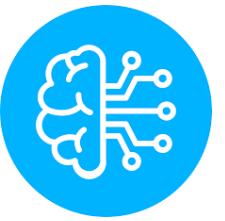
Instead of training a deep network from scratch for your task:

- Take a network trained on a different domain for a different **source task**
- Adapt it for your domain and your **target task**

Variations:

- Same domain, different task
- Different domain, same task





Transfer Learning

Example: Suppose you are a camera surveillance company, hired by the government to install cameras in the city to identify parked cars.

Goal: Identify cars that have overstayed the parking limit and deserve a ticket, during 9 am to 5pm

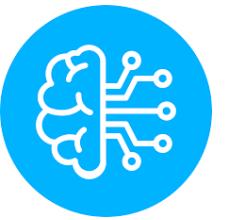
Solution: Implement a neural network that takes camera footage as input and identify cars (parked in the day light)

Due to the success of the application, the government want to extend the time period to night also:

Option 1: start from scratch and build a new models trained on both images during day and night

Option 2: use your pretrained model, adjust it on night images=> additional training on night images





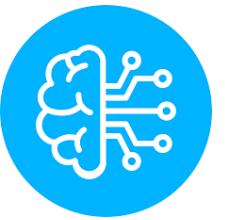
Transfer learning

Scenario: Improving Skin Cancer Detection

Imagine a medical research team working on a project to improve the early detection of skin cancer, a critical and life-saving task. The team faces a challenge: they have a limited dataset of skin lesion images from their local hospital, which isn't sufficient for training a robust deep learning model.

Problem: The team needs to create a skin cancer detection system that can accurately classify skin lesions as malignant or benign, but they lack the necessary labeled data.

Solution: Transfer Learning overcome data limitations and improve the accuracy of their skin cancer detection model.



Transfer learning

Scenario: Personalized Healthcare with Limited Medical Data

Imagine a small rural clinic in a remote area that lacks access to advanced medical expertise and technology. The clinic serves a diverse population with a range of medical conditions, but they have limited patient data for making accurate diagnoses and treatment decisions.

Problem: The clinic's healthcare practitioners face challenges in providing personalized medical care due to the scarcity of patient data and limited access to specialized medical knowledge.

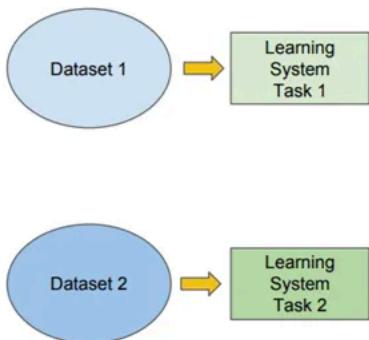
Solution: Transfer Learning => empowers a rural clinic to provide personalized healthcare services despite data limitations.

Transfer Learning



Traditional ML

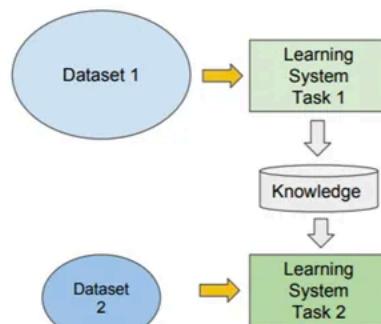
- Isolated, single task learning:
 - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks



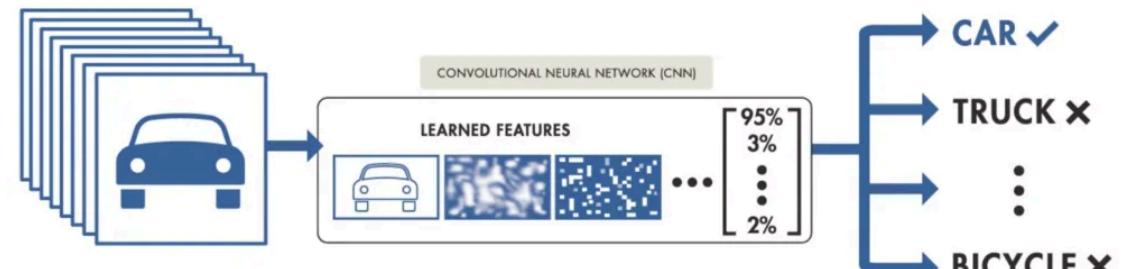
vs

Transfer Learning

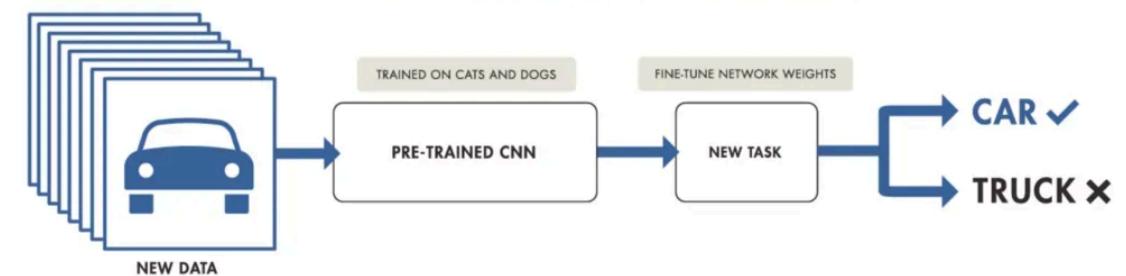
- Learning of a new tasks relies on the previous learned tasks:
 - Learning process can be faster, more accurate and/or need less training data

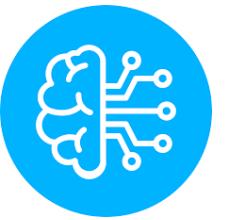


TRAINING FROM SCRATCH



TRANSFER LEARNING

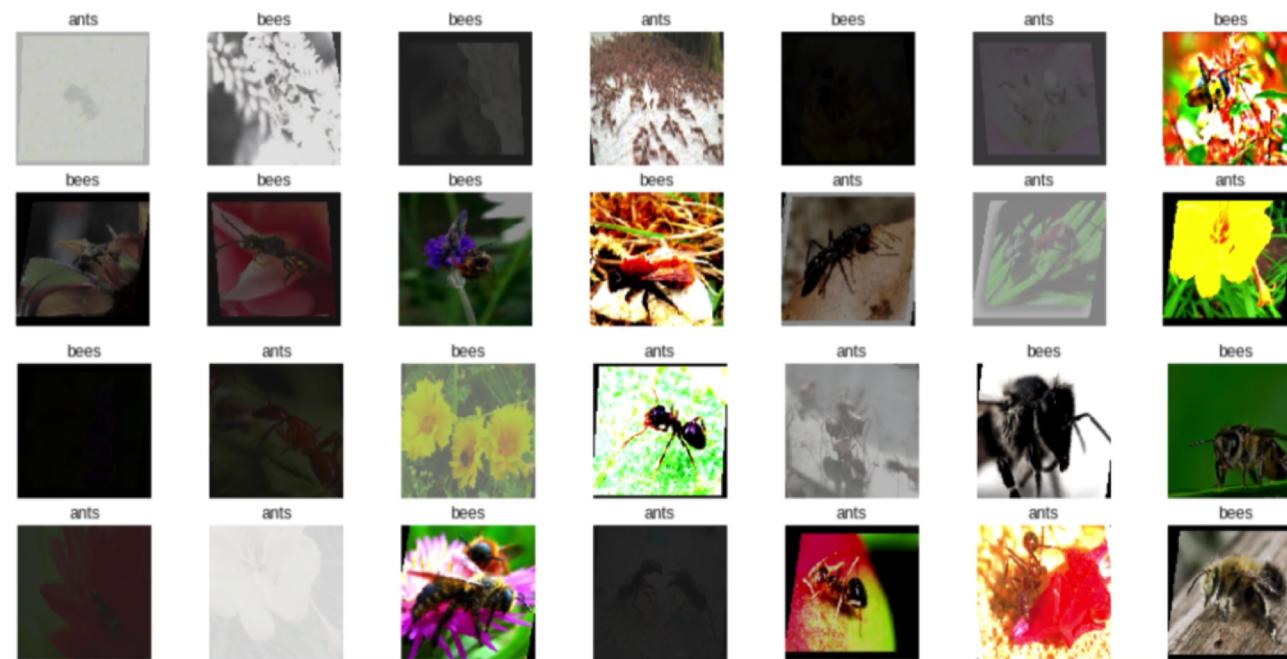


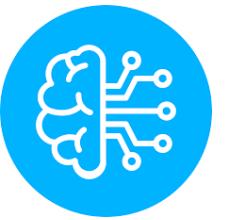


Transfer Learning

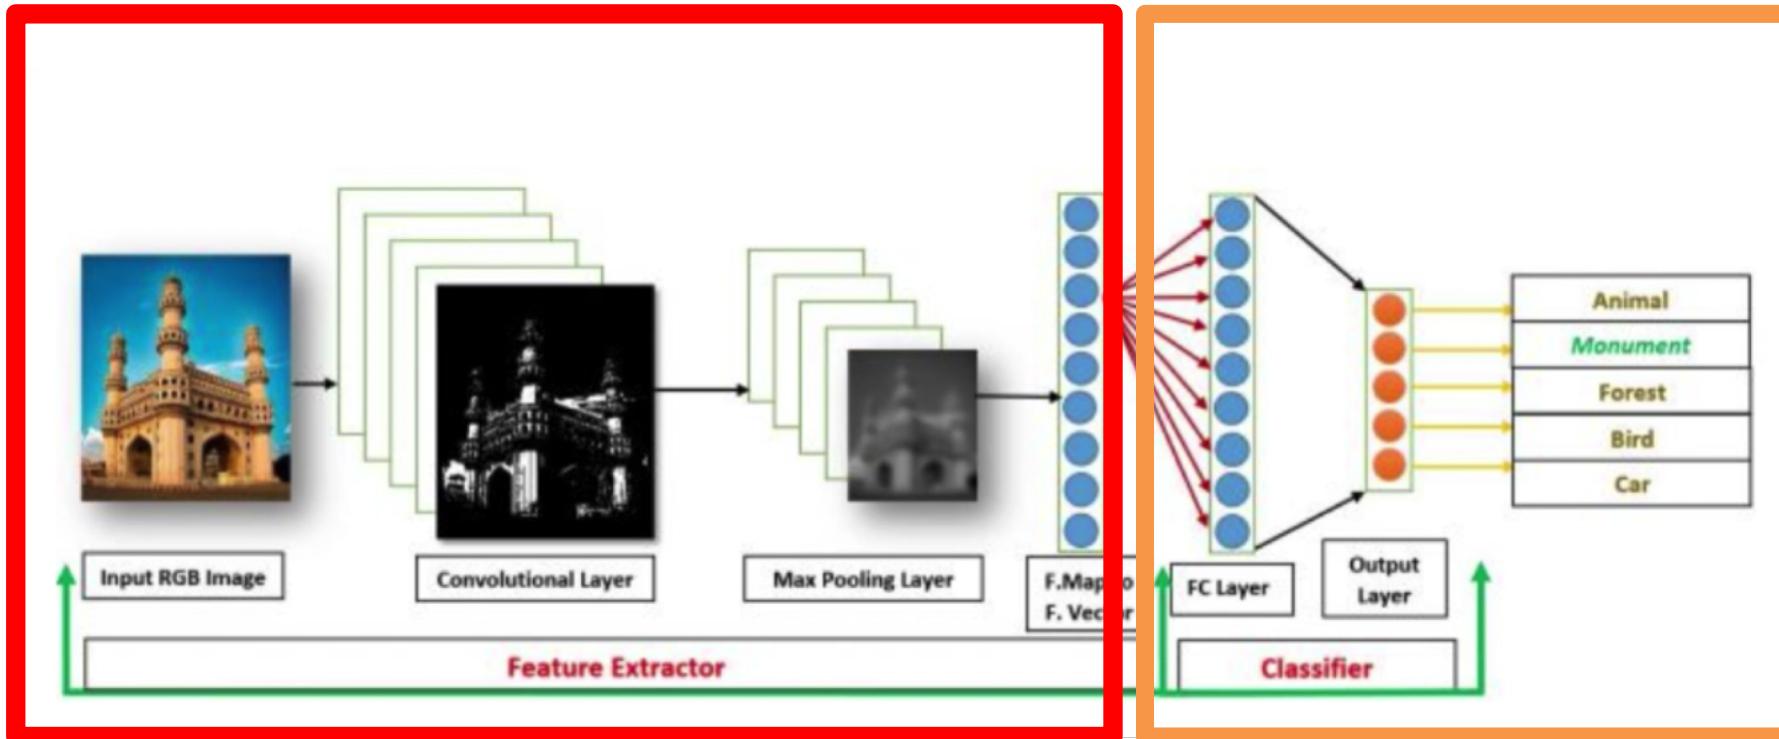
Ant/bee classification

- Not enough labelled data (400 images)
- Effective pretrained model already exists





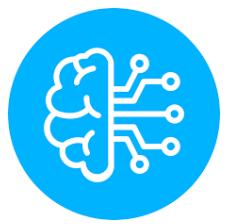
Transfer Learning



Freeze

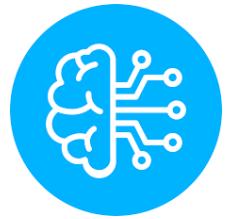
Re-train

Transfer Learning



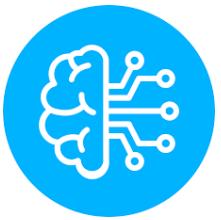
When:

- **Limited Data Availability:** Instead of training a deep learning model from scratch, you can start with a pre-trained model on a large dataset and fine-tune it on your smaller dataset. This leverages the knowledge already captured in the pre-trained model.
- **Similar Tasks:** Transfer learning is most effective when the source task (the one the pre-trained model was trained on) is related to the target task (your specific problem).
- **Reduced Training Time:** Can significantly reduce training time because you're building on pre-existing knowledge.
- **Resource Constraints:** If you don't have access to large-scale computing resources, transfer learning can be a more practical option than training a model from scratch.



Transfer Learning

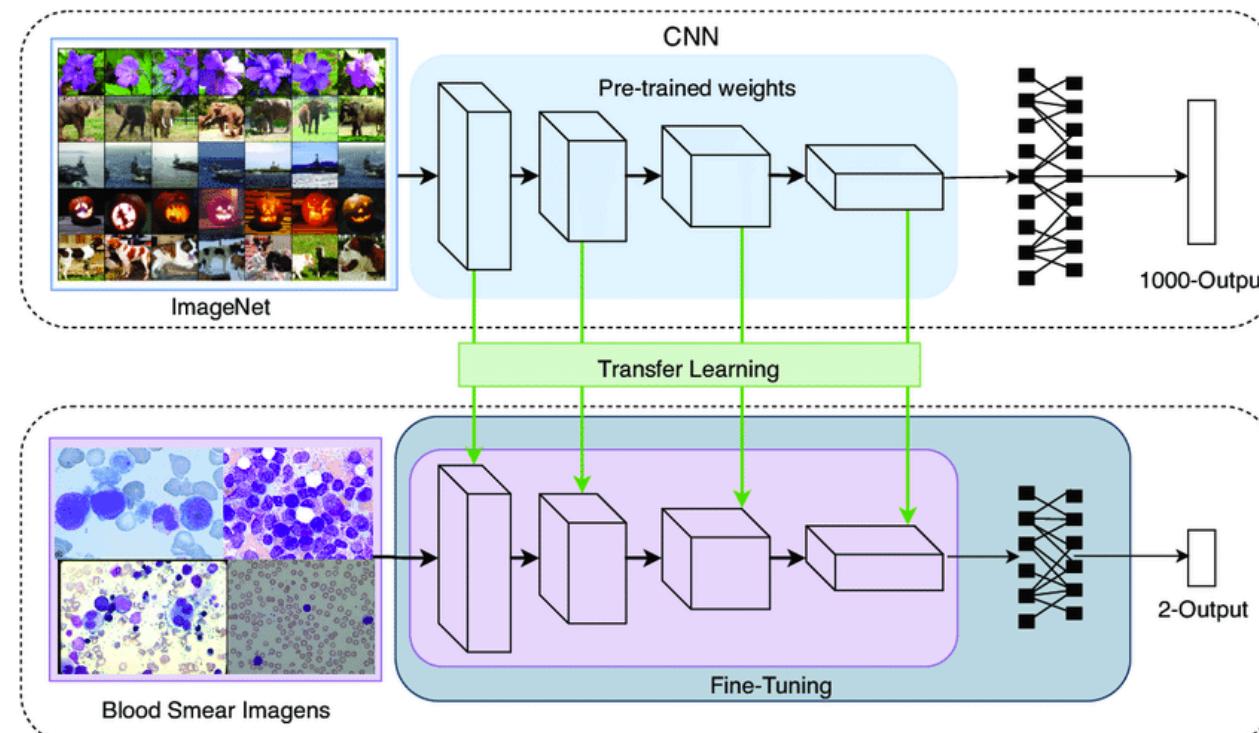
- Broad concept that encompasses various techniques for utilizing knowledge from a pre-trained model in a new task.
- Involves taking a model that has been trained on a **source task** and applying it to a **related target task**. The primary idea is to transfer knowledge learned from the source task to the target task, which can help **improve** performance or **reduce the amount of data** required for training
- Typically replaces the final layers (the task-specific output layers) of the pre-trained model with new layers that match the target task. These new layers are randomly initialized.
- The weights of the pre-trained model are often **frozen**, meaning they are not updated during training on the target task. Only the weights of the new layers are trained on the target data.
- **Transfer learning is commonly used when the source and target tasks are related but not identical, and when the target dataset is relatively small.**

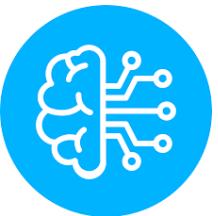


Transfer Learning

Fine tuning

- Fine-tuning is a technique within transfer learning where not only the final layers but also some earlier layers of the pre-trained model are modified to adapt it to the target task.
- While transfer learning freezes all the pre-trained layers and only trains the new layers, fine-tuning goes a step further by allowing the pre-trained layers to be updated





Transfer Learning

Fine tuning

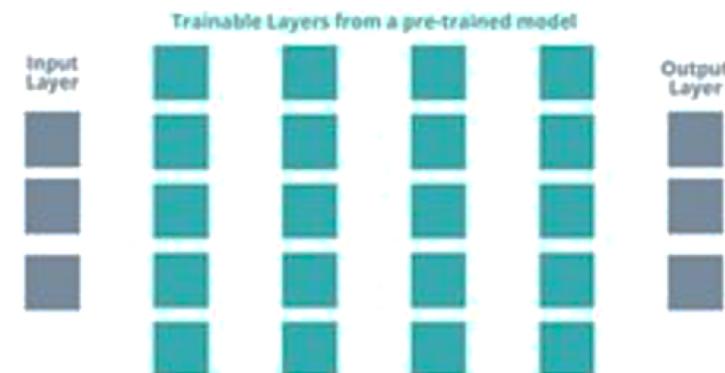
Transfer Learning: How Feature Extraction & Fine-Tuning work?

Feature Extraction

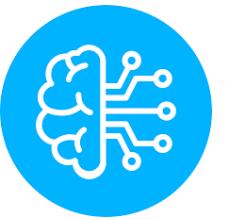


In feature extraction, you **freeze the pre-trained model** layers to preserve existing learning and **add new layers** to learn additional information.

Fine Tuning



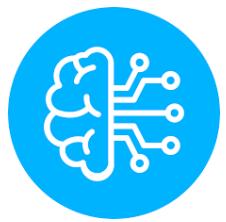
In fine-tuning, you **unfreeze the entire model** and **train it with a lower learning rate** to adapt to new challenges.



Transfer learning

Scenario	Use Fine-Tuning	Freeze Layers (Feature Extraction)
Similar source and target tasks	Yes	Sometimes
Moderate to sufficient target data	Yes	Sometimes (especially with limited data)
Task requires task-specific features	Yes	Sometimes (if not needed)
Complex target task	Yes	Sometimes (depends on task complexity)
Limited target data	Sometimes (if enough data)	Yes
Dissimilar source and target tasks	Sometimes (if some features transfer)	Yes
Limited computational resources	Sometimes (if resources are available)	Yes
Privacy and security concerns with the pre-trained model	Sometimes (if no sensitive data)	Yes

Transfer Learning

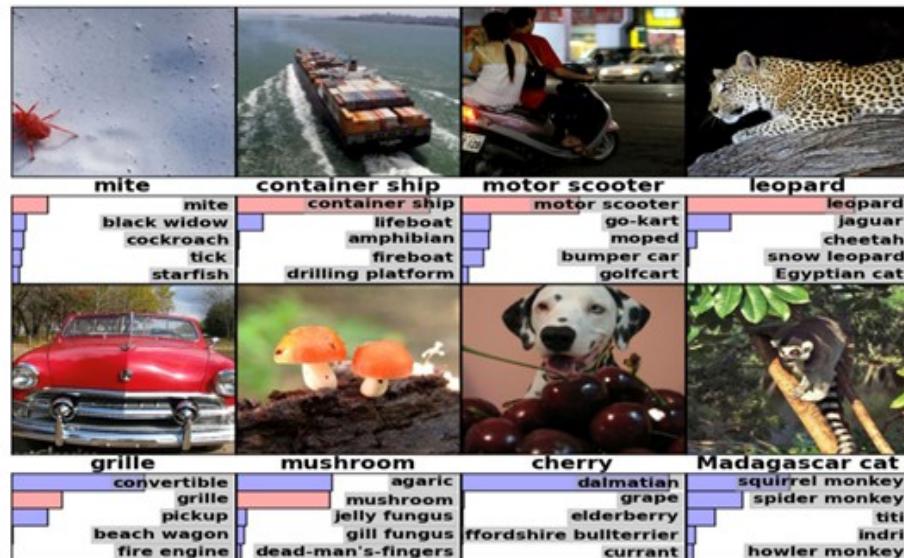


ImageNet Large Scale Visual Recognition Challenge (ILSVRC), was an annual computer vision competition that aimed to evaluate the performance of computer vision algorithms on a large-scale image dataset. The challenge was initiated in 2010 and ran until 2017.

ImageNet Challenge



- 1,000 object classes (categories).
- Images:
 - 1.2 M train
 - 100k test.





Transfer Learning

Main task: image classification.

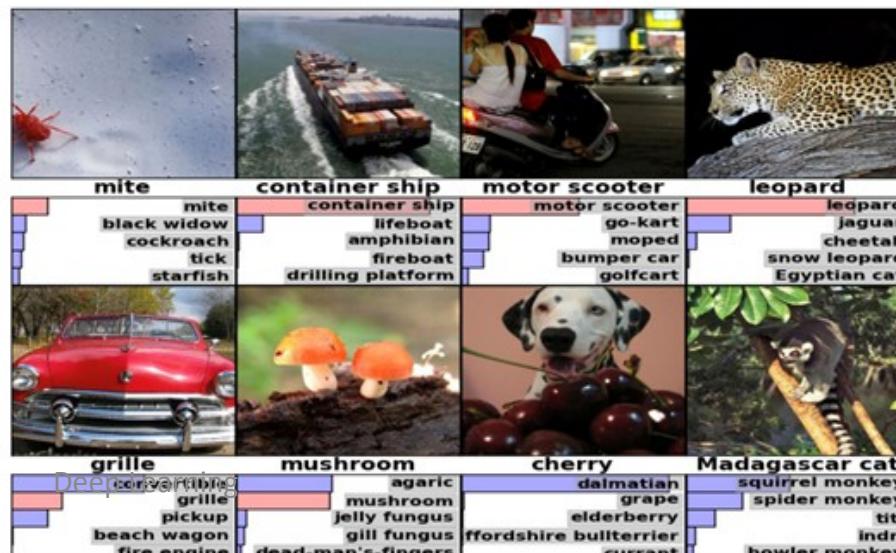
Participants were provided with a dataset called ImageNet, which contained millions of labeled images across thousands of categories.

The goal was to develop algorithms that could accurately classify images into the correct categories. Initially, it focused on 1,000 object categories.

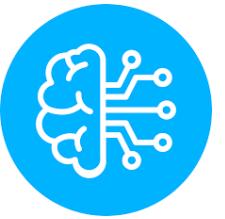
ImageNet Challenge



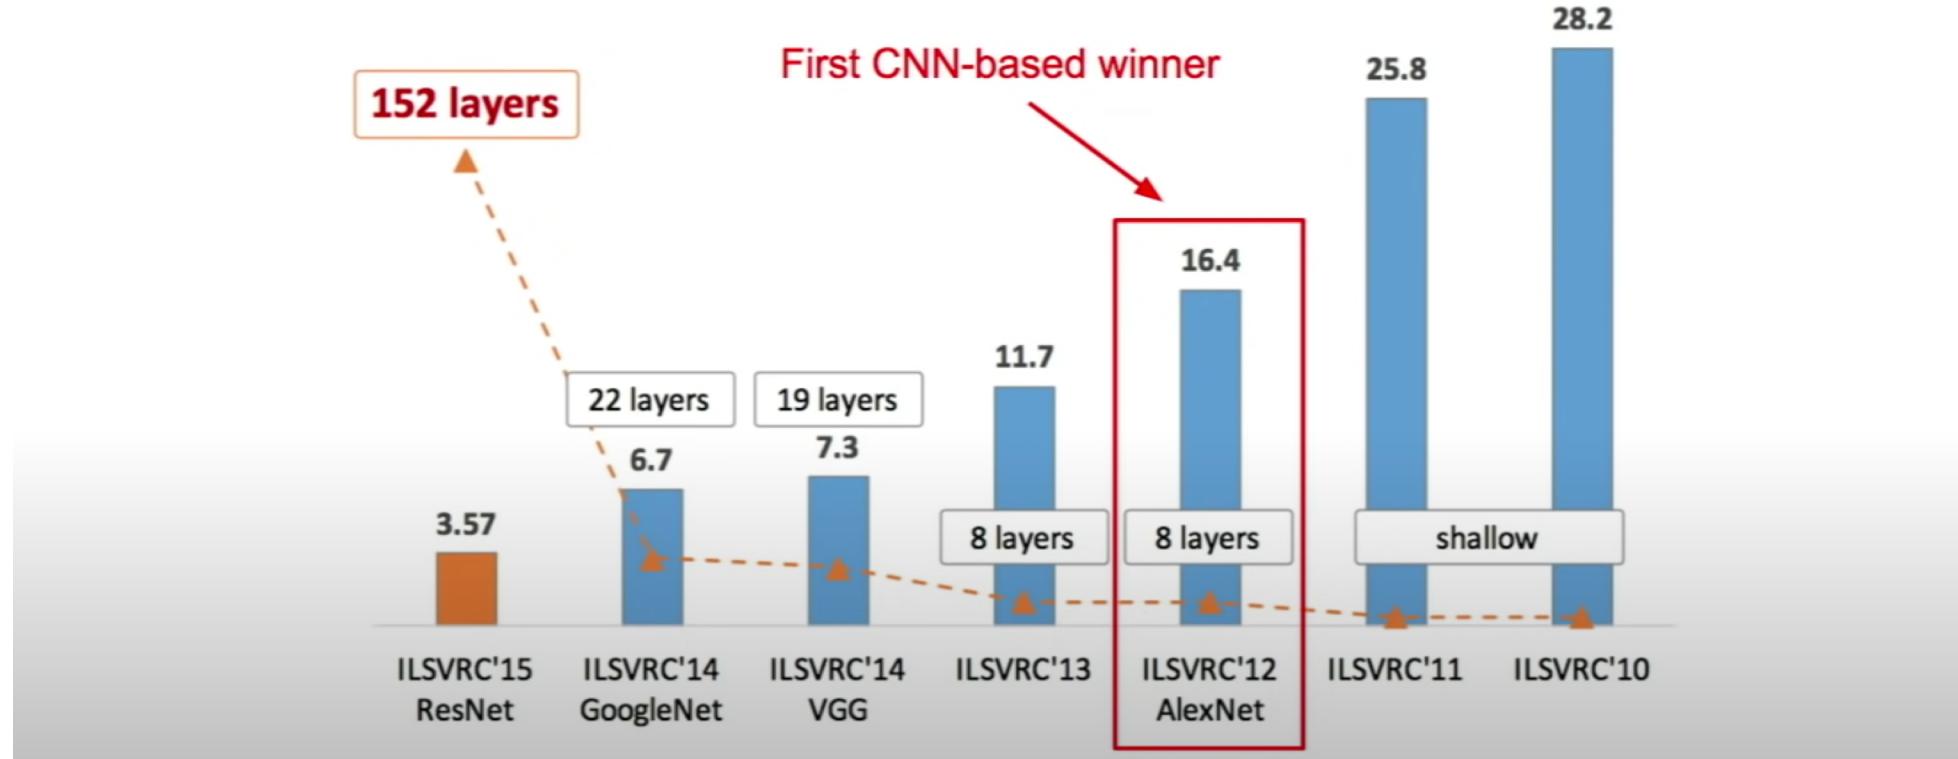
- 1,000 object classes (categories).
- Images:
 - 1.2 M train
 - 100k test.



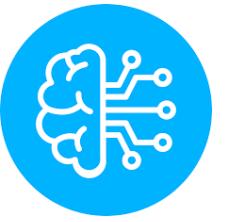
ImageNet Challenge



ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



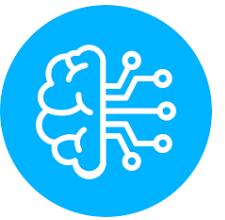
ImageNet



In the ImageNet Large Scale Visual Recognition Challenge (ILSVRC), "top-1" and "top-5" errors are metrics used to evaluate the performance of image classification models.

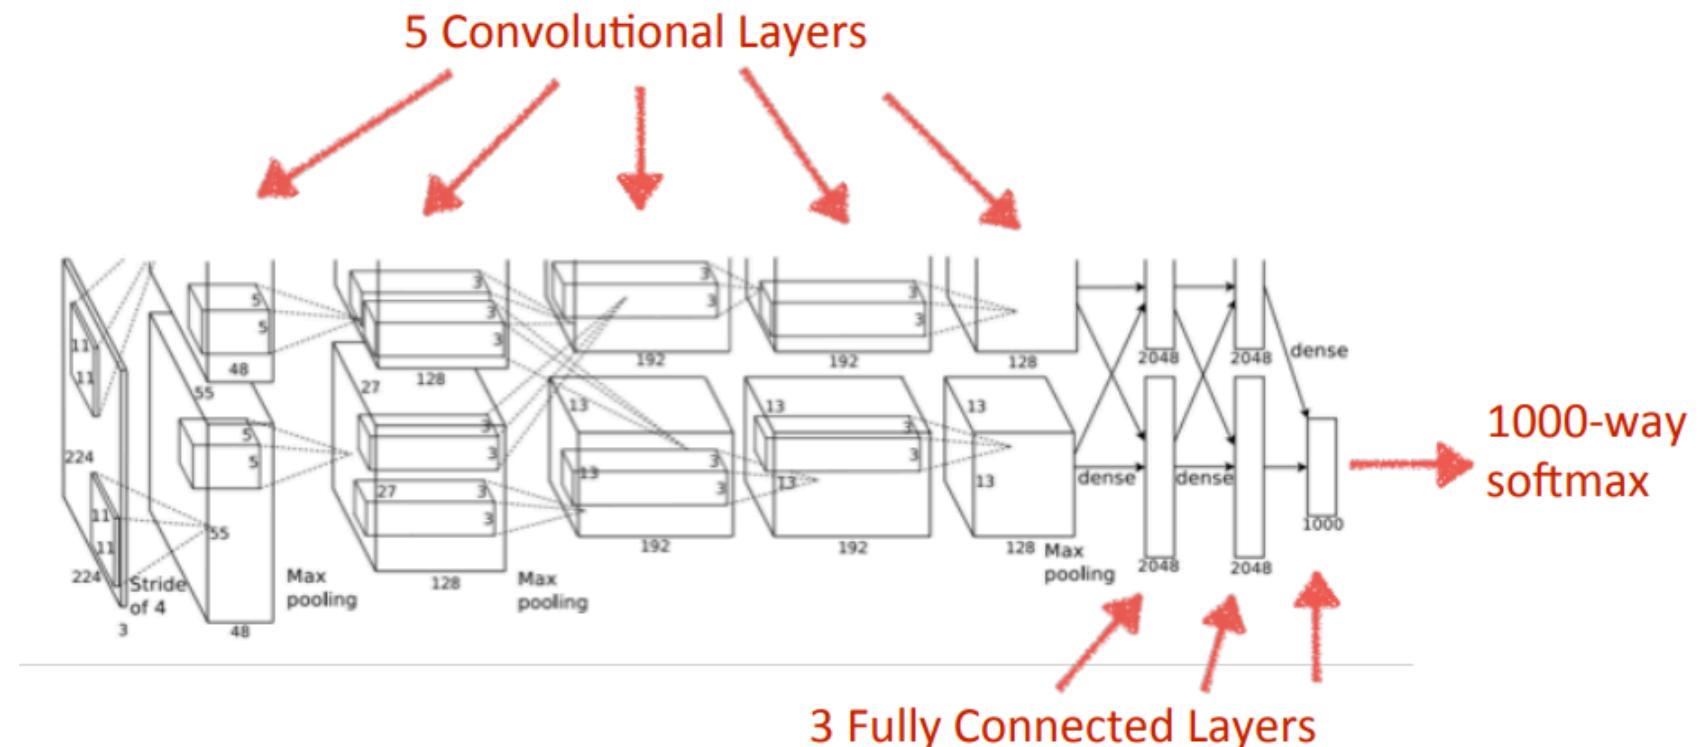
Top-1 Error: Measures the percentage of test images for which the correct category (the single most likely category predicted by the model) is not the same as the ground truth label. In other words, it calculates the proportion of images where the model's top prediction is incorrect.

Top-5 Error: Measures the percentage of test images for which the correct category is not among the top five most likely categories predicted by the model. In other words, it calculates the proportion of images where the correct label is not found within the model's top five predictions.

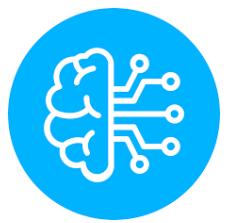


AlexNet

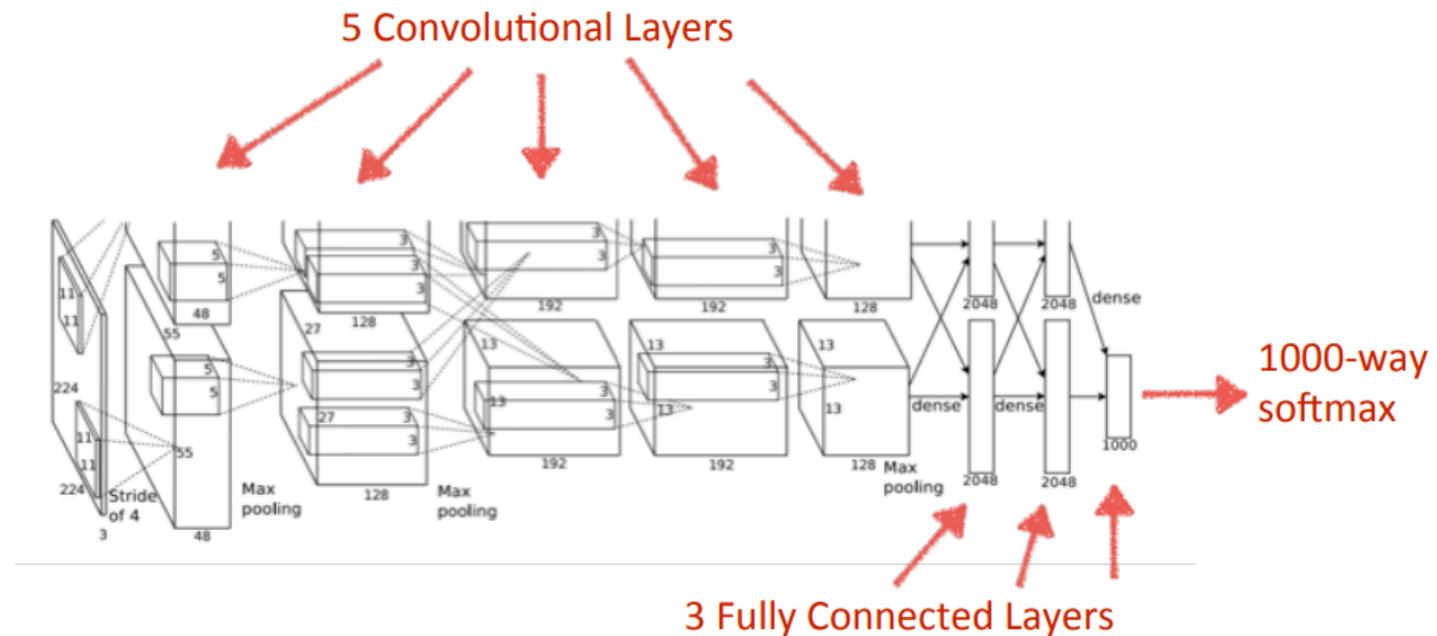
- 2012: one of the first deep convolutional neural networks (CNNs) to demonstrate the effectiveness of deep architectures for image classification.
- Top 5 test error rate of 15.4% (2nd **26,2 %**)



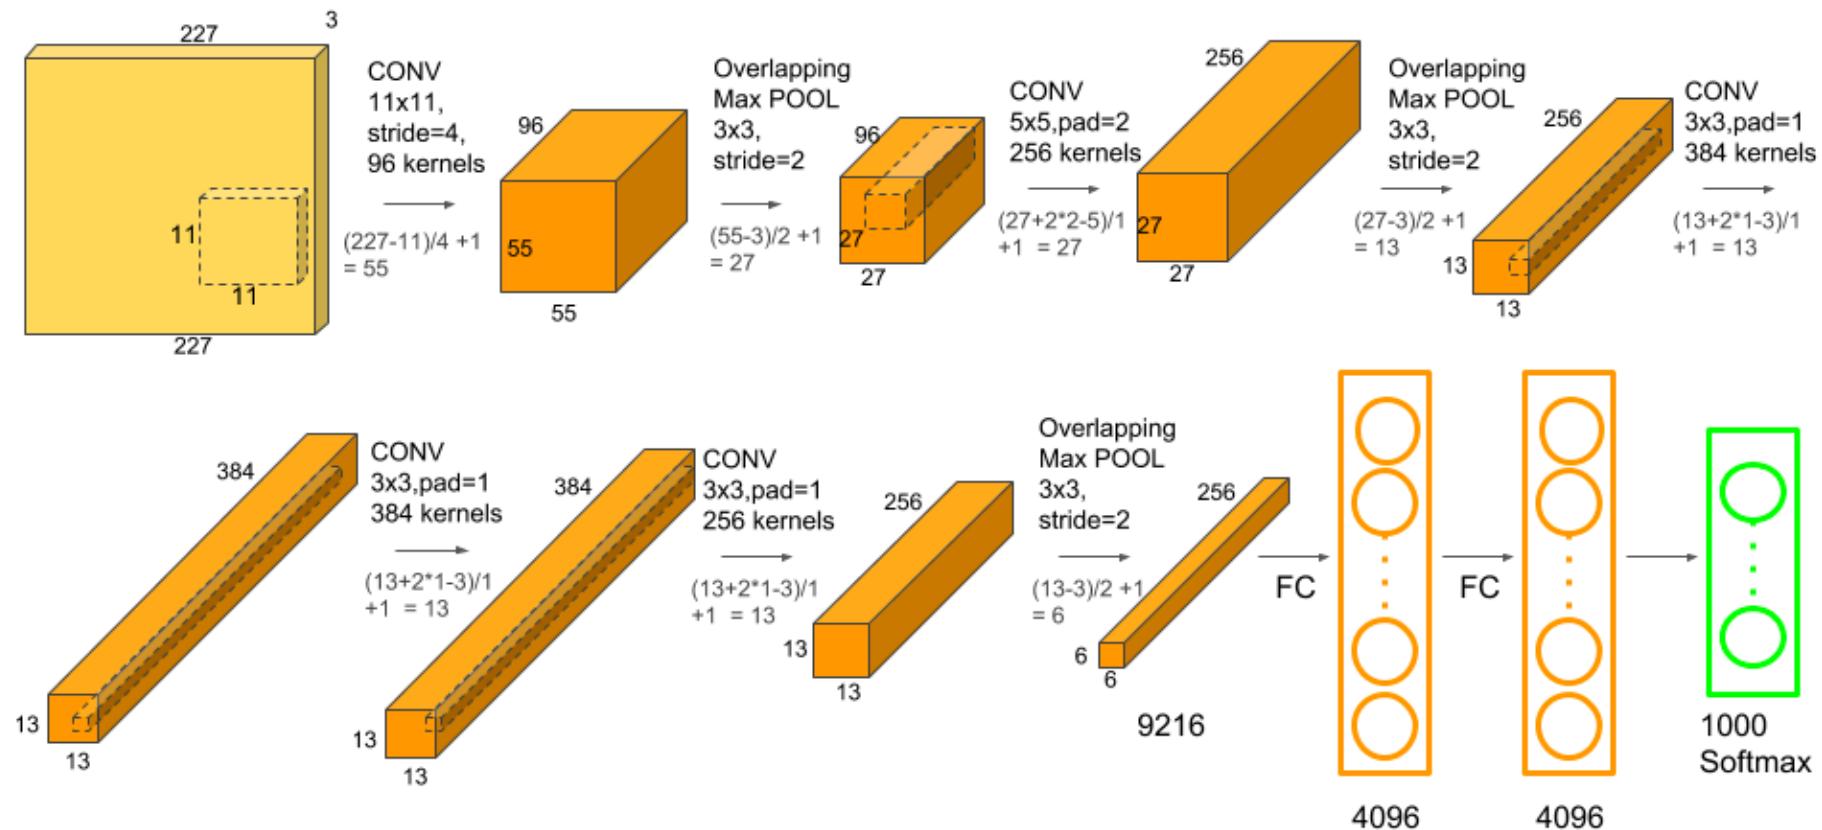
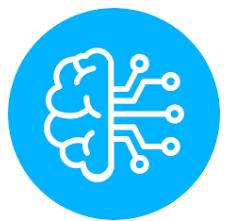
AlexNet



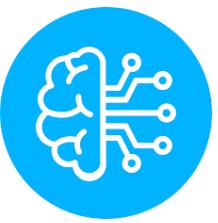
- Used ReLU for non-linearity
- Used data augmentation techniques
- Implemented dropout layers to prevent overfitting
- Trained on two GTX 580 GPUs for five to six days



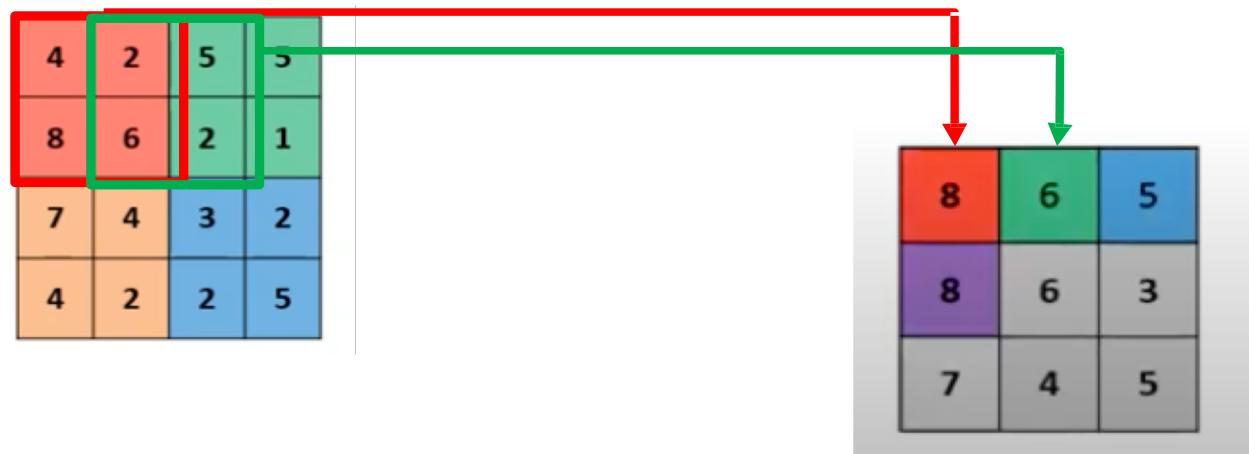
AlexNet



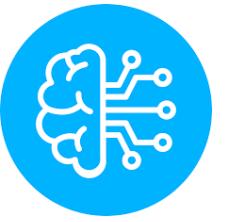
AlexNet



Overlapping pooling:

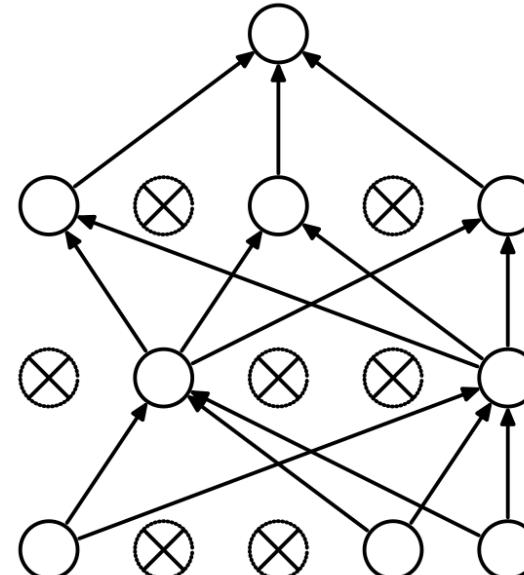
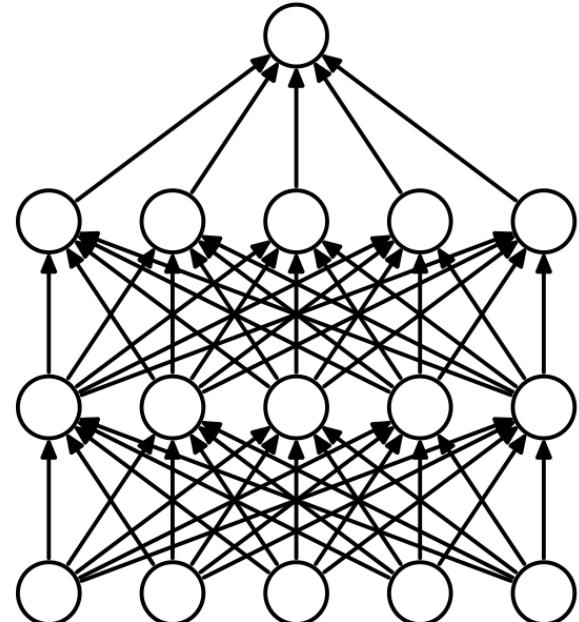


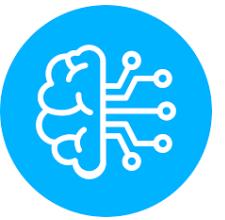
AlexNet



Dropout

- Dropout is a regularization technique commonly used in deep learning, to prevent overfitting.
- Overfitting occurs when a model performs well on the training data but fails to generalize effectively to unseen data.





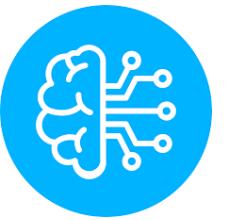
Transfer learning

Dropout

During Training: dropout randomly **deactivates** (sets to zero) a fraction of neurons in a neural network layer on each forward and backward pass. These deactivated neurons do not contribute to the computations for that particular iteration. The fraction of neurons that are deactivated is a **hyperparameter**, typically set between 0.2 and 0.5.

Random Deactivation: The key idea behind dropout is that it prevents individual neurons from becoming overly specialized or dependent on specific input features.

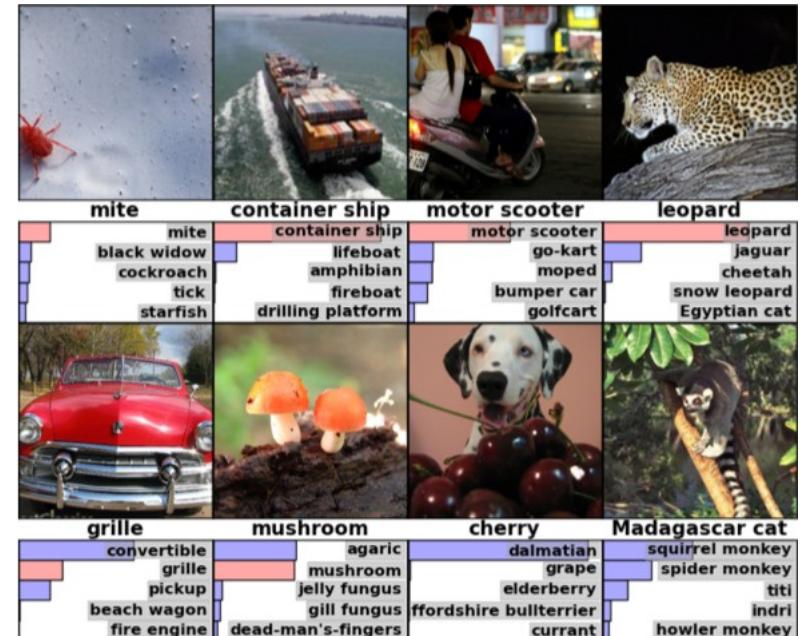
Testing Phase: dropout is turned off, and all neurons are active. The predictions are made using the full network.



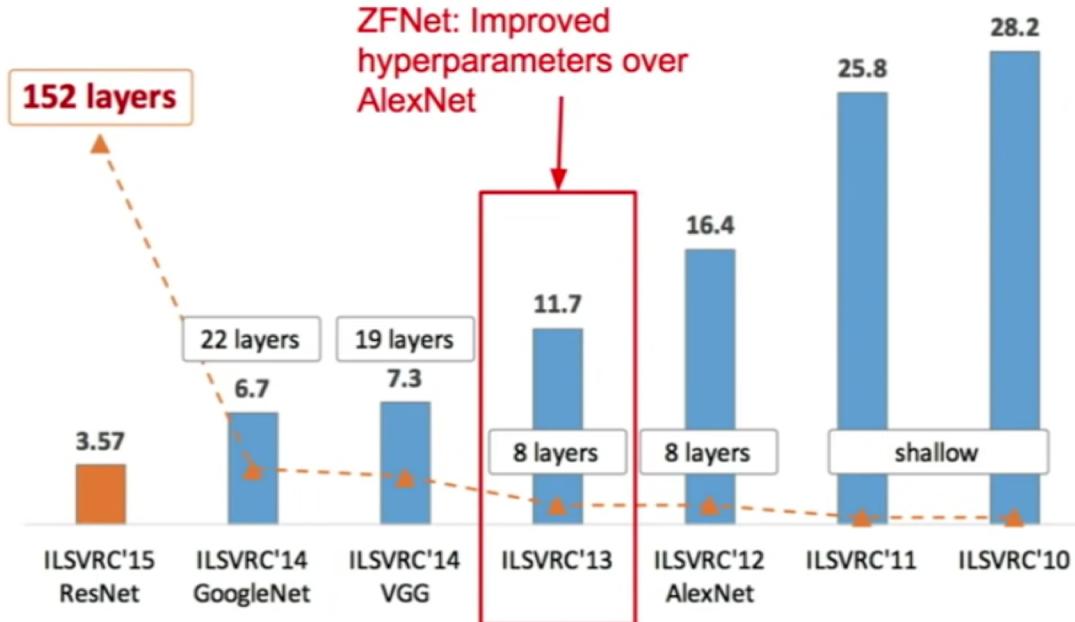
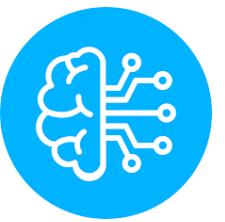
AlexNet

The Results.

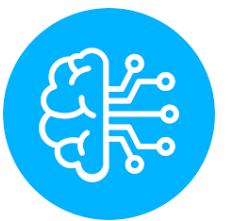
- On the 2010 version of the ImageNet competition, the best model achieved **47.1% top-1 error and 28.2% top-5 error.**
- AlexNet vastly outpaced this with a **37.5% top-1 error and a 17.0% top-5 error.**
- AlexNet **won the 2012 ImageNet competition** with a top-5 error rate of 15.3%, compared to the second place top-5 error rate of 26.2%.



ZFNet

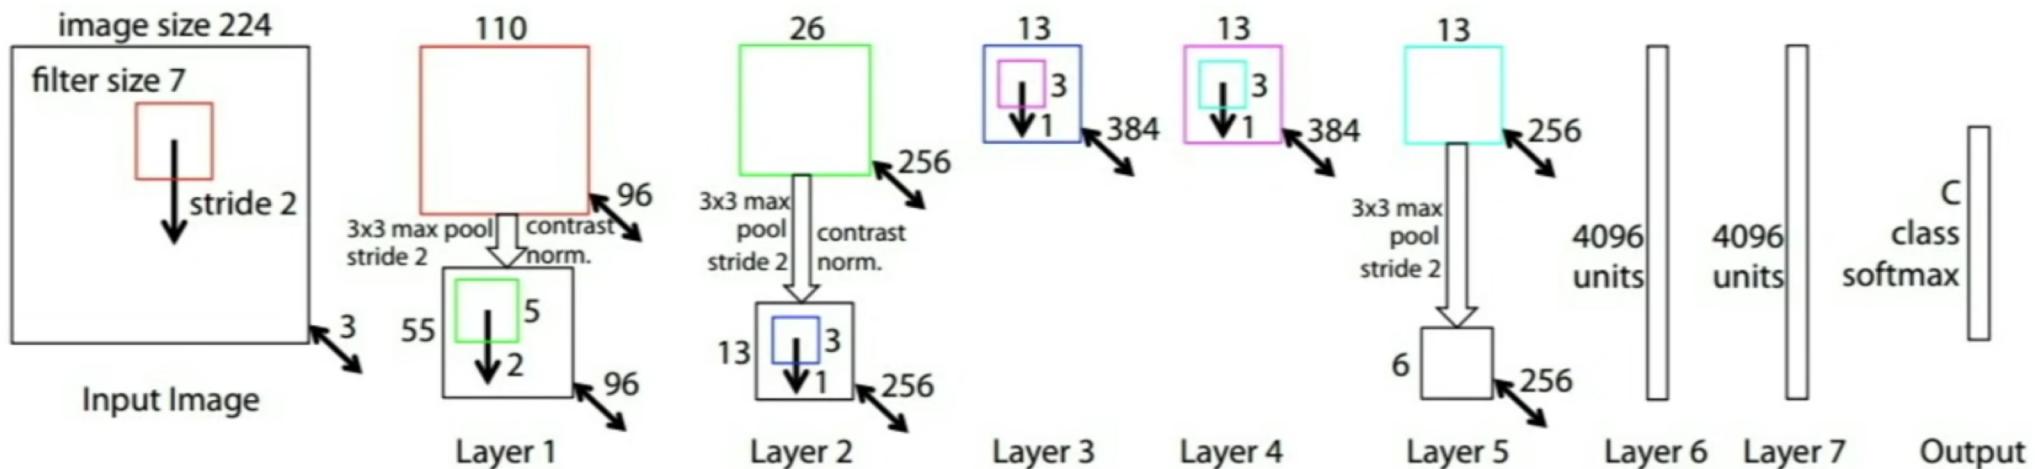


ZFNet



ZFNet

[Zeiler and Fergus, 2013]

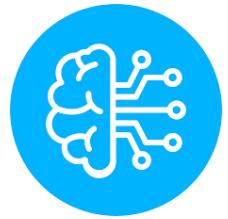


AlexNet but:

CONV1: change from (11x11 stride 4) to (7x7 stride 2)

CONV3,4,5: instead of 384, 384, 256 filters use 512, 1024, 512

TODO: remake figure



Parameters

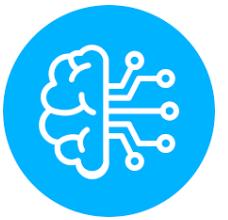
How to calculate parameters for a CNN?

```
1 my_input = Input(shape = (10, 10, 1))
2 my_output = Convolution2D(1, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 model = Model(my_input, my_output)

1 model.summary()
```

- Black and white image
- 10*10 image
- 1 convolutional filter
- 3*3 filter
- The output => 8*8 image

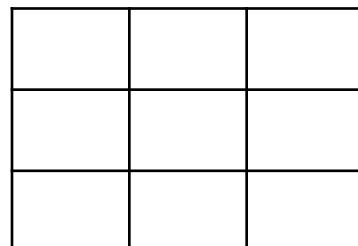
Layer (type)	Output Shape	Param #
<hr/>		
input_17 (InputLayer)	[(None, 10, 10, 1)]	0
conv2d_19 (Conv2D)	(None, 8, 8, 1)	10
activation_19 (Activation)	(None, 8, 8, 1)	0
<hr/>		
Total params:	10	
Trainable params:	10	
Non-trainable params:	0	



Parameters

```
1 my_input = Input(shape = (10, 10, 1))
2 my_output = Convolution2D(1, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 model = Model(my_input, my_output)
```

```
1 model.summary()
```

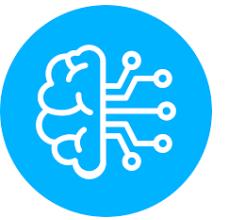


3*3 filter + bias = 10

- Black and white image (Grayscale)
- 10*10 image
- 1 convolutional filter (3*3)
- The output => 8*8 image

Layer (type)	Output Shape	Param #
input_17 (InputLayer)	[(None, 10, 10, 1)]	0
conv2d_19 (Conv2D)	(None, 8, 8, 1)	10
activation_19 (Activation)	(None, 8, 8, 1)	0

Total params: 10
Trainable params: 10
Non-trainable params: 0



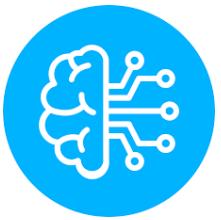
Parameters

Example 2

```
1 my_input = Input(shape = (10, 10, 1))
2 my_output = Convolution2D(5, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(2, 3)(my_output)
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```

- Black and white image (Grayscale)
- 10*10 image
- 5 convolutional filters (3*3)

Layer (type)	Output Shape	Param #
=====		
input_19 (InputLayer)	[None, 10, 10, 1]	0
conv2d_22 (Conv2D)	(None, 8, 8, 5)	50
activation_22 (Activation)	(None, 8, 8, 5)	0
conv2d_23 (Conv2D)	(None, 6, 6, 2)	92
activation_23 (Activation)	(None, 6, 6, 2)	0
=====		
Total params: 142		
Trainable params: 142		
Non-trainable params: 0		



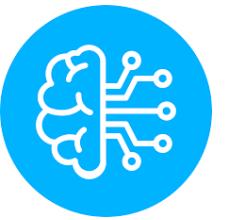
Parameters

Example 2

```
1 my_input = Input(shape = (10, 10, 1))
2 my_output = Convolution2D(5, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(2, 3)(my_output)
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```

- Black and white image (Grayscale)
- 10*10 image
- 5 convolutional filters (3*3)
- Output?

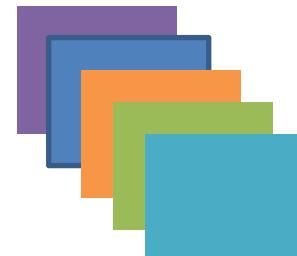
Layer (type)	Output Shape	Param #
=====		
input_19 (InputLayer)	[None, 10, 10, 1]	0
conv2d_22 (Conv2D)	(None, 8, 8, 5)	50
activation_22 (Activation)	(None, 8, 8, 5)	0
conv2d_23 (Conv2D)	(None, 6, 6, 2)	92
activation_23 (Activation)	(None, 6, 6, 2)	0
=====		
Total params: 142		
Trainable params: 142		
Non-trainable params: 0		



Parameters

Example 2

```
1 my_input = Input(shape = (10, 10, 1))
2 my_output = Convolution2D(5, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(2, 3)(my_output)
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```



10*10 image

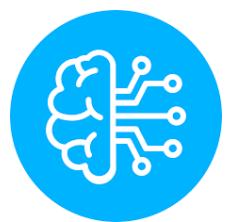
5 filters (3*3)

5 images (8*8)

- Black and white image (Grayscale)
- 10*10 image
- 5 convolutional filters (3*3)
- Output?

layer (type)	Output Shape	Param #
input_19 (InputLayer)	[None, 10, 10, 1]	0
conv2d_22 (Conv2D)	(None, 8, 8, 5)	50
activation_22 (Activation)	(None, 8, 8, 5)	0
conv2d_23 (Conv2D)	(None, 6, 6, 2)	92
activation_23 (Activation)	(None, 6, 6, 2)	0
Total params:	142	
Trainable params:	142	
Non-trainable params:	0	

Each filter has 10 parameters associated => 50 parameters



Parameters

```
1 my_input = Input(shape = (10, 10, 1))
2 my_output = Convolution2D(5, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(2, 3)(my_output) highlighted
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```

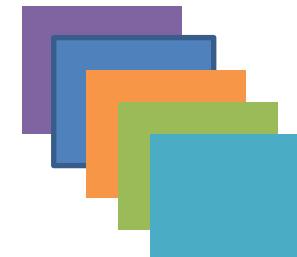
Input

- Five 8*8 images (8*8*5)
- 2 convolutional filters (3*3)

Output?

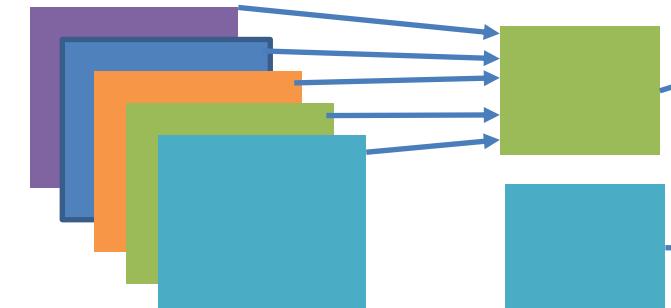
Layer (type)	Output Shape	Param #
input_19 (InputLayer)	[None, 10, 10, 1]	0
conv2d_22 (Conv2D)	(None, 8, 8, 5)	50
activation_22 (Activation)	(None, 8, 8, 5)	0
conv2d_23 (Conv2D)	(None, 6, 6, 2)	92
activation_23 (Activation)	(None, 6, 6, 2)	0

Total params: 142
Trainable params: 142
Non-trainable params: 0



10*10 image

5 filters (3*3)



5 images (8*8)

Each filter has 5 channels
(3,3,5)



2 Filters (3*3)



(6,6,2)



Parameters

```
1 my_input = Input(shape = (10, 10, 1))
2 my_output = Convolution2D(5, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(2, 3)(my_output)
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```

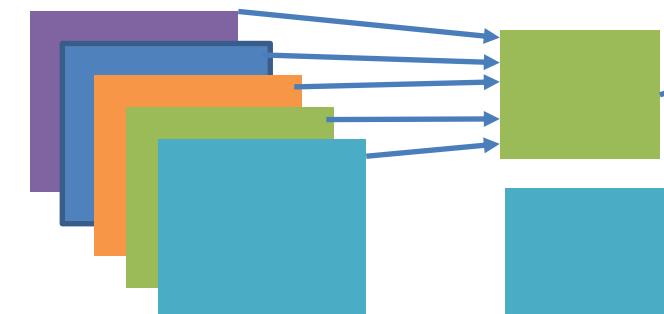
Input

- Five 8*8 images (8*8*5)
- 2 convolutional filters (3*3)

Output?



5 filters (3*3)



5 images (8*8)
Deep Learning

2 Filters (3*3)



(6,6,2)

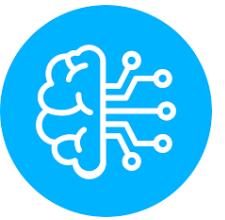
Each filter has 5 channels
(3,3,5)

6*6 image
6*6 image
(6,6,2)

For each filter:
 $3 \times 3 \times 5 = 45$
We add the bias => 46

Two filters= 92

Layer (type)	Output Shape	Param #
input_19 (InputLayer)	[None, 10, 10, 1]	0
conv2d_22 (Conv2D)	(None, 8, 8, 5)	50
activation_22 (Activation)	(None, 8, 8, 5)	0
conv2d_23 (Conv2D)	(None, 6, 6, 2)	92
activation_23 (Activation)	(None, 6, 6, 2)	0
=====		
Total params: 142		
Trainable params: 142		
Non-trainable params: 0		



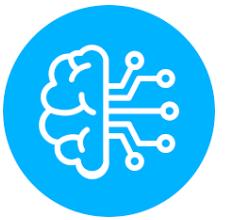
Parameters

```
1 my_input = Input(shape = (100, 100, 3))
2 my_output = Convolution2D(8, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(1, 3)(my_output)
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```

Output?

- image of 100*100*3 (colored image)
 - 8 convolutional filters of 3*3
- ⇒ 98*98*3 (number of filters)

⇒ How much parameters we have?



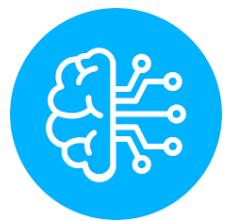
Parameters

```
1 my_input = Input(shape = (100, 100, 3))
2 my_output = Convolution2D(8, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(1, 3)(my_output)
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```

Output?

- image of 100*100*3 (colored image)
- 8 convolutional filters of 3*3
⇒ 98*98*3 (number of filters)

⇒ How much parameters we have?
⇒ The filters adapt to the image depth => the filters are 3*3*3 filters
⇒ 27 parameters + bias
⇒ 28
⇒ 28*8 => 224



Parameters

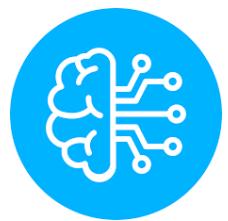
```
1 my_input = Input(shape = (100, 100, 3))
2 my_output = Convolution2D(8, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(1, 3)(my_output)
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```

Output?

- image of 100*100*3 (colored image)
 - 8 convolutional filters of 3*3
⇒ 98*98*3 (number of filters)
- ⇒ How much parameters we have?
⇒ The filters matchs the image depth => 3*3*3 filters
⇒ 27 parameters + bias
⇒ 28
⇒ 28*8 => 224

Layer (type)	Output Shape	Param #
input_20 (InputLayer)	[(None, 100, 100, 3)]	0
conv2d_24 (Conv2D)	(None, 98, 98, 8)	224
activation_24 (Activation)	(None, 98, 98, 8)	0
conv2d_25 (Conv2D)	(None, 96, 96, 1)	73
activation_25 (Activation)	(None, 96, 96, 1)	0

Total params: 297		
Trainable params: 297		
Non-trainable params: 0		



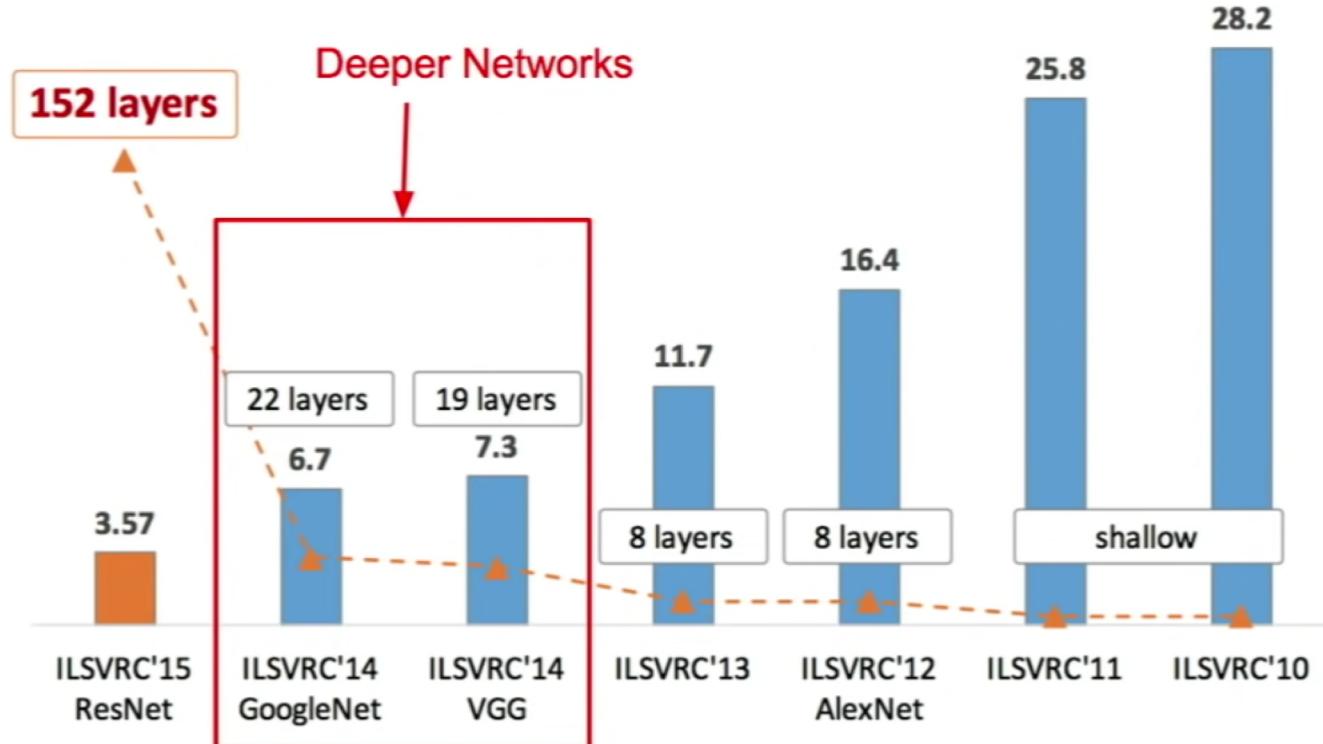
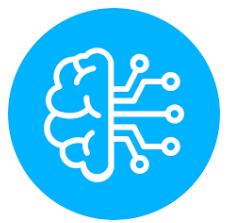
Parameters

```
1 my_input = Input(shape = (100, 100, 3))
2 my_output = Convolution2D(8, 3)(my_input)
3 my_output = Activation('sigmoid')(my_output)
4 my_output = Convolution2D(1, 3)(my_output)  
5 my_output = Activation('sigmoid')(my_output)
6 model = Model(my_input, my_output)
```

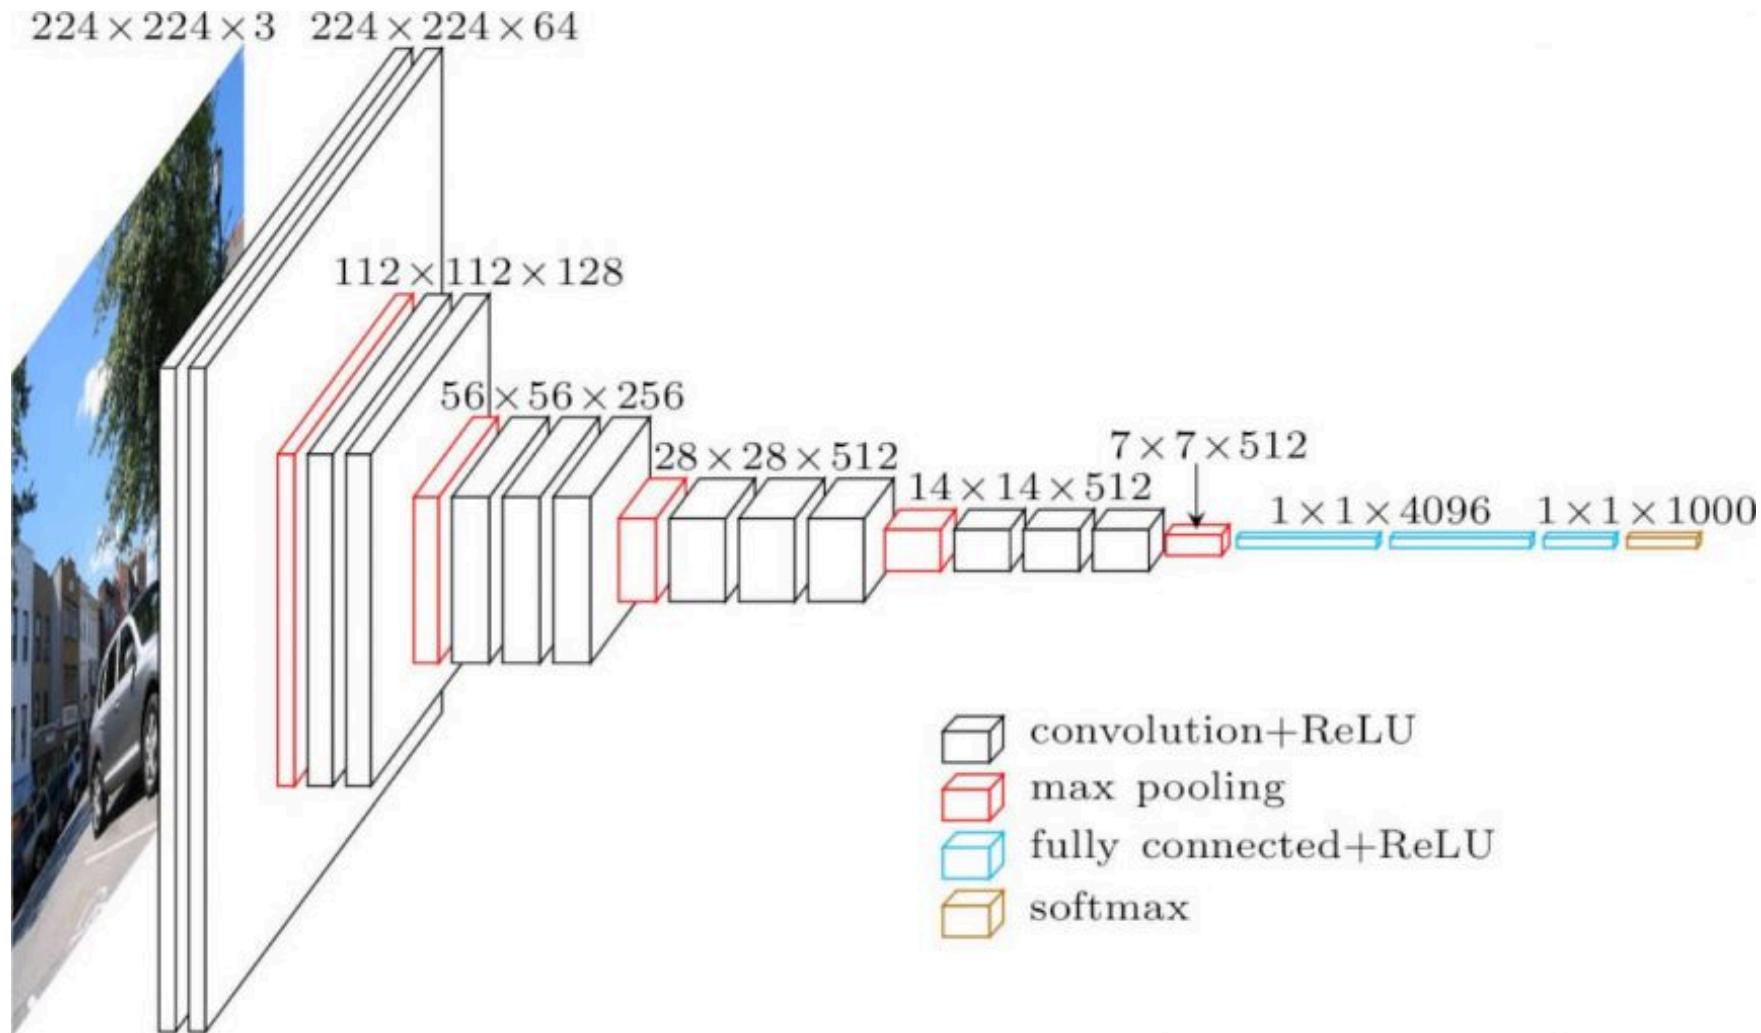
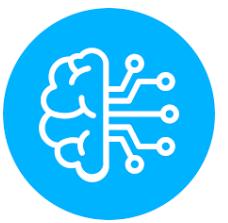
How did we get 73?

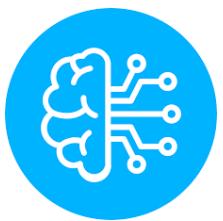
Layer (type)	Output Shape	Param #
input_20 (InputLayer)	[(None, 100, 100, 3)]	0
conv2d_24 (Conv2D)	(None, 98, 98, 8)	224
activation_24 (Activation)	(None, 98, 98, 8)	0
conv2d_25 (Conv2D)	(None, 96, 96, 1)	73
activation_25 (Activation)	(None, 96, 96, 1)	0
<hr/>		
Total params: 297		
Trainable params: 297		
Non-trainable params: 0		

Deeper Networks



VGG





Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Small filters, Deeper networks

8 layers (AlexNet)

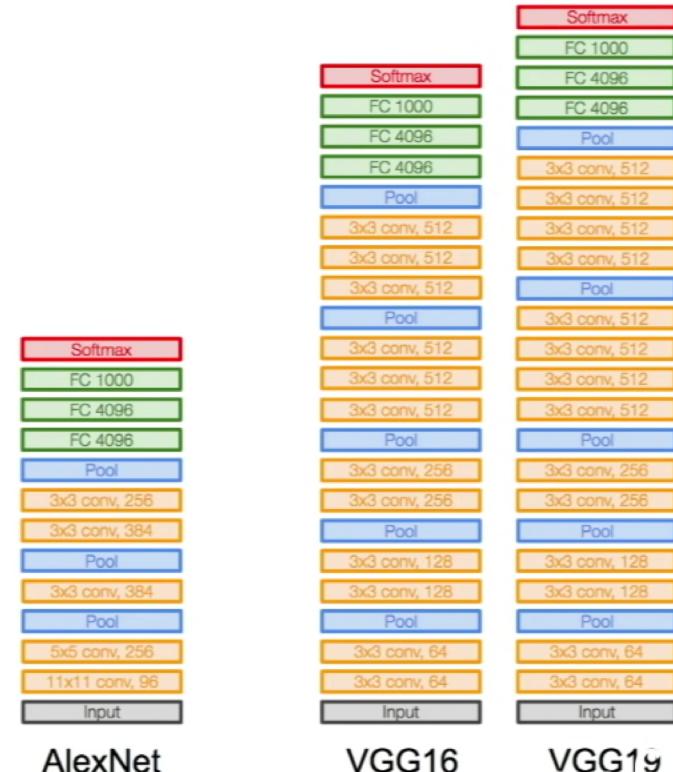
-> 16 - 19 layers (VGG16Net)

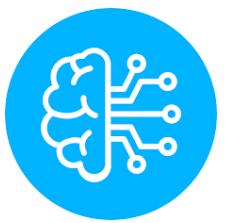
Only 3x3 CONV stride 1, pad 1
and 2x2 MAX POOL stride 2

11.7% top 5 error in ILSVRC'13

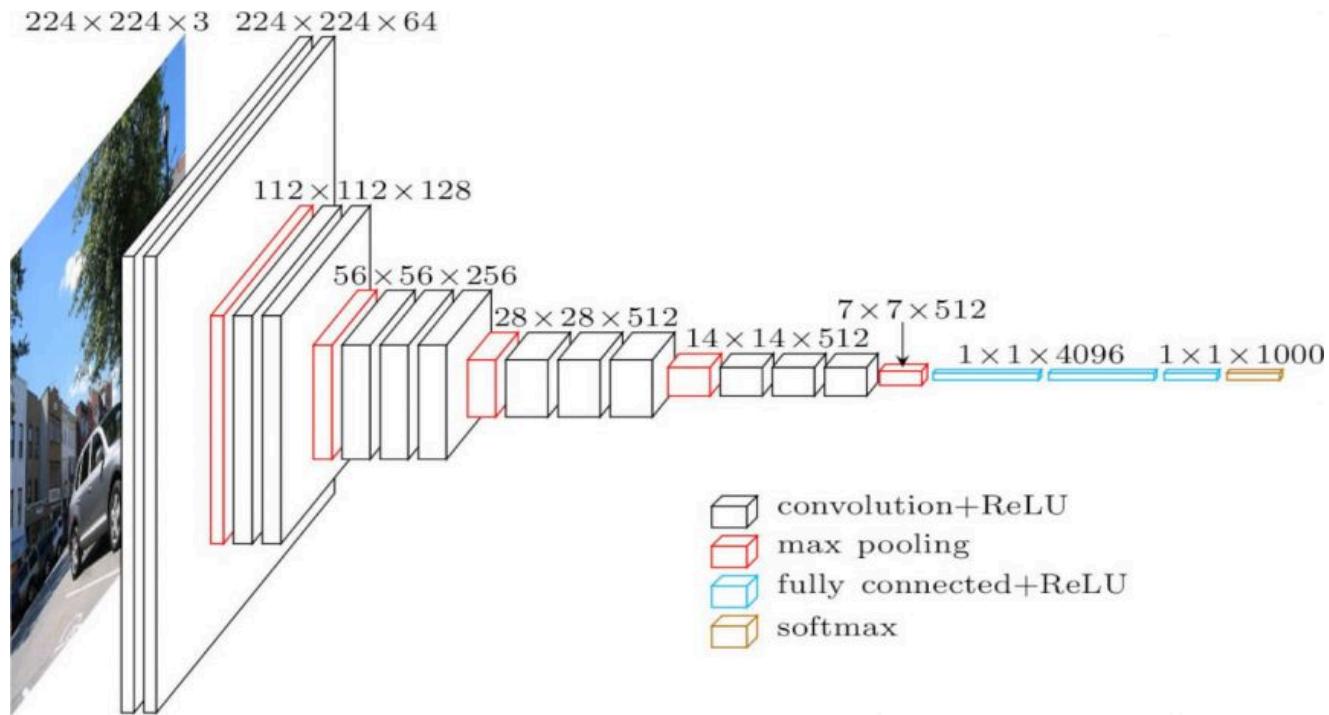
(ZFNet)

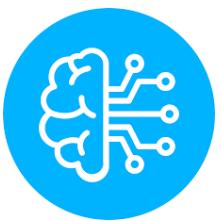
-> 7.3% top 5 error in ILSVRC'14





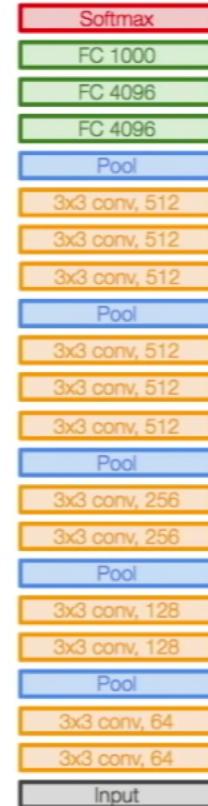
- Simplicity and depth (16/19 layers)
- 7.3% error rate
- The use of only 3×3 sized filters
- As the spatial size of the input volumes at each layer increases, the depth of the volume increases (increased number of filters)
- Trained on 4 Nvidia Titan BlackGPUs for Two weeks



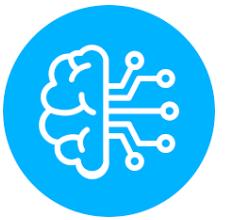


INPUT: [224x224x3] memory: $224 \times 224 \times 3 = 150\text{K}$ params: 0 (not counting biases)
 CONV3-64: [224x224x64] memory: $224 \times 224 \times 64 = 3.2\text{M}$ params: $(3 \times 3 \times 3) \times 64 = 1,728$
 CONV3-64: [224x224x64] memory: $224 \times 224 \times 64 = 3.2\text{M}$ params: $(3 \times 3 \times 64) \times 64 = 36,864$
 POOL2: [112x112x64] memory: $112 \times 112 \times 64 = 800\text{K}$ params: 0
 CONV3-128: [112x112x128] memory: $112 \times 112 \times 128 = 1.6\text{M}$ params: $(3 \times 3 \times 64) \times 128 = 73,728$
 CONV3-128: [112x112x128] memory: $112 \times 112 \times 128 = 1.6\text{M}$ params: $(3 \times 3 \times 128) \times 128 = 147,456$
 POOL2: [56x56x128] memory: $56 \times 56 \times 128 = 400\text{K}$ params: 0
 CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800\text{K}$ params: $(3 \times 3 \times 128) \times 256 = 294,912$
 CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800\text{K}$ params: $(3 \times 3 \times 256) \times 256 = 589,824$
 CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800\text{K}$ params: $(3 \times 3 \times 256) \times 256 = 589,824$
 POOL2: [28x28x256] memory: $28 \times 28 \times 256 = 200\text{K}$ params: 0
 CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400\text{K}$ params: $(3 \times 3 \times 256) \times 512 = 1,179,648$
 CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
 CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
 POOL2: [14x14x512] memory: $14 \times 14 \times 512 = 100\text{K}$ params: 0
 CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
 CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
 CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
 POOL2: [7x7x512] memory: $7 \times 7 \times 512 = 25\text{K}$ params: 0
 FC: [1x1x4096] memory: 4096 params: $7 \times 7 \times 512 \times 4096 = 102,760,448$
 FC: [1x1x4096] memory: 4096 params: $4096 \times 4096 = 16,777,216$
 FC: [1x1x1000] memory: 1000 params: $4096 \times 1000 = 4,096,000$

TOTAL memory: $24\text{M} * 4 \text{ bytes} \approx 96\text{MB} / \text{image}$ (only forward! ~ 2 for bwd)
TOTAL params: 138M parameters

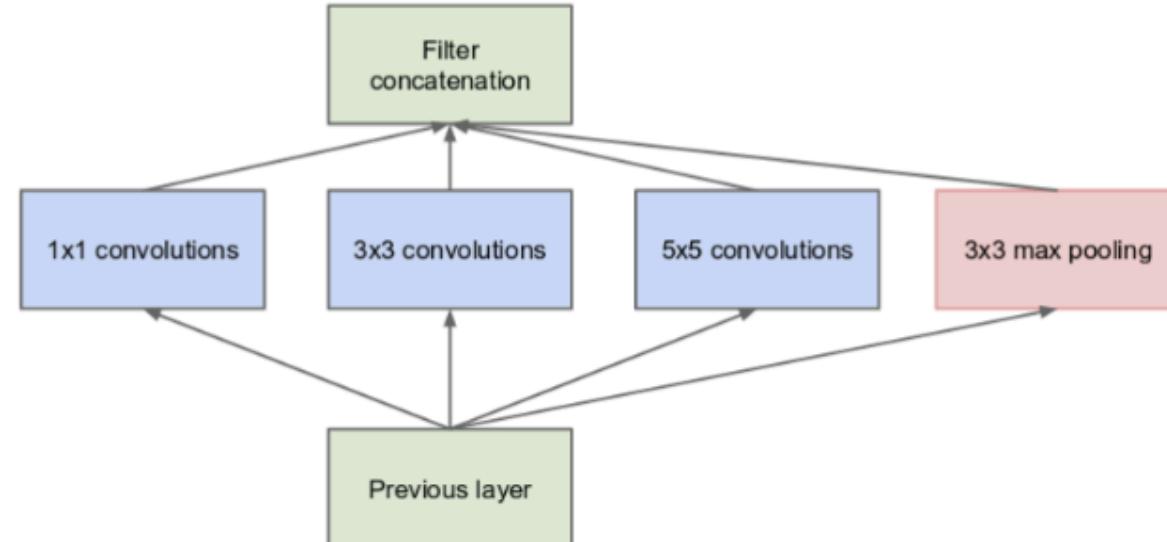


VGG16

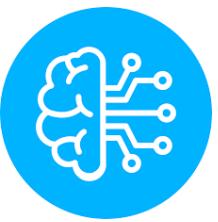


GoogleNet

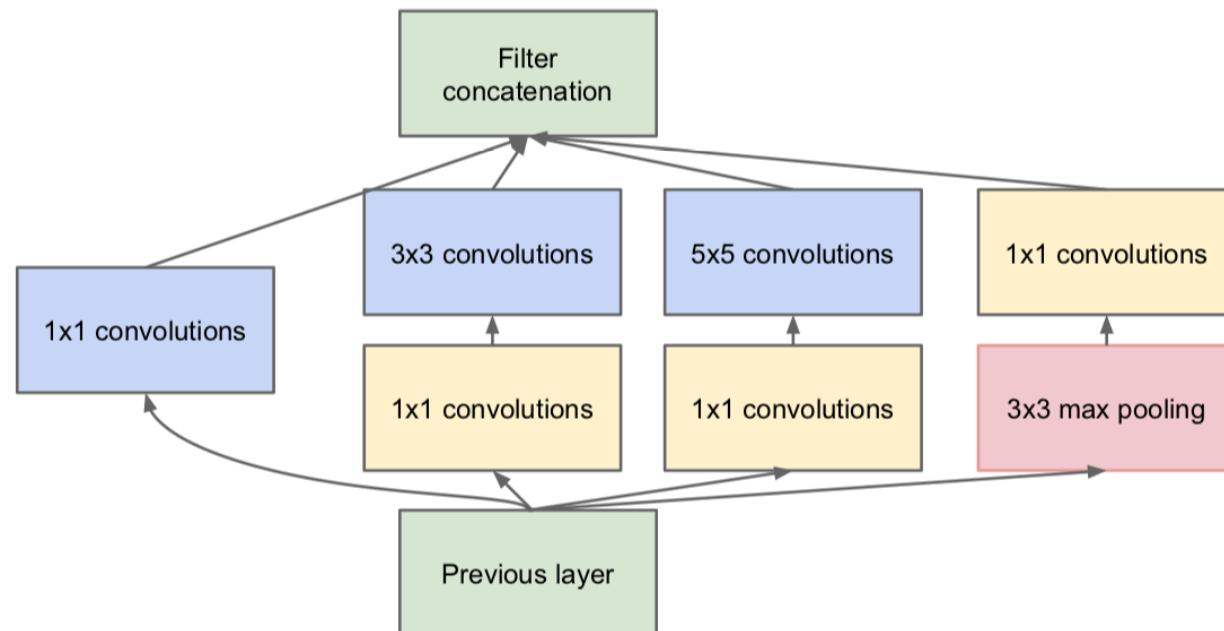
- Top 5 error rate of 6.7%
- Key point: parallel convolutions with output concatenation
- 100 layers in total
- No use of fully connected layers
- Uses 12x fewer parameters than AlexNet (**only 5 million parameters**)



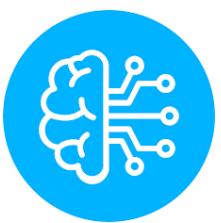
GoogleNet



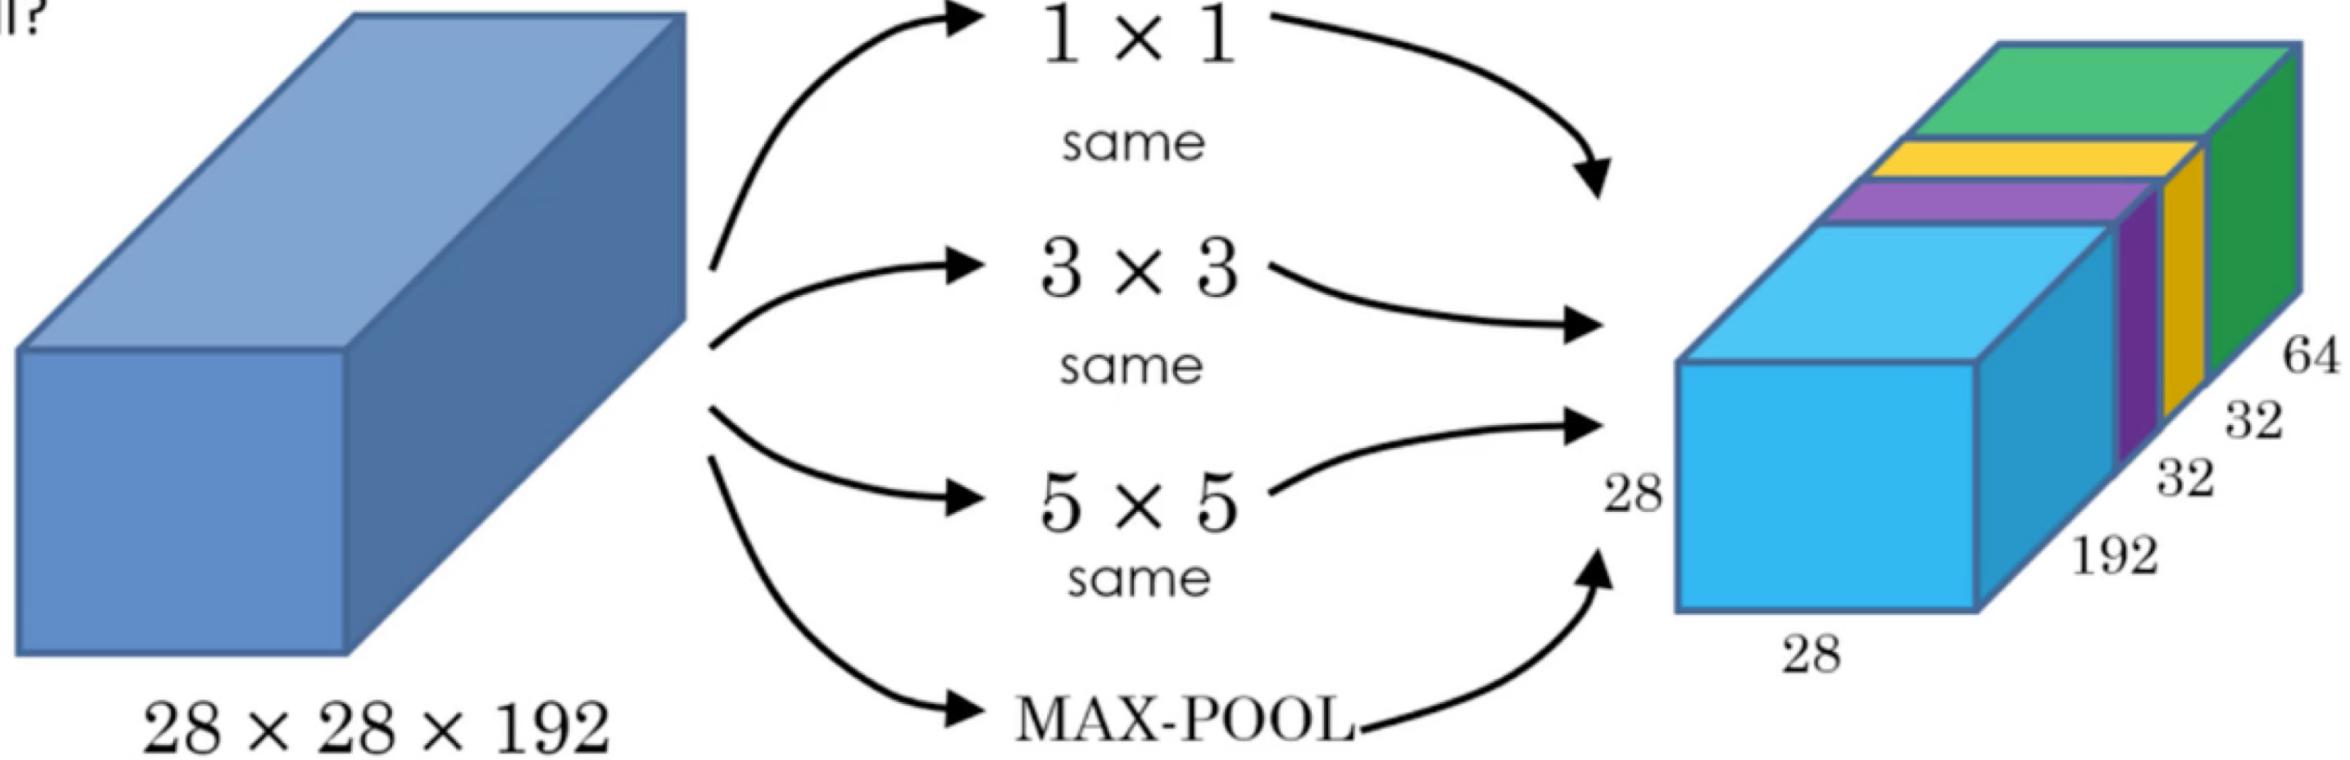
Do you want to use 1×1 filter, or 3×3 , or 5×5 ? Use a pooling layer? Or use them all together?



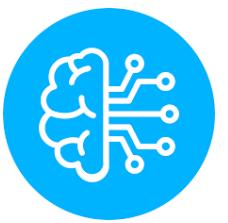
GoogleNet



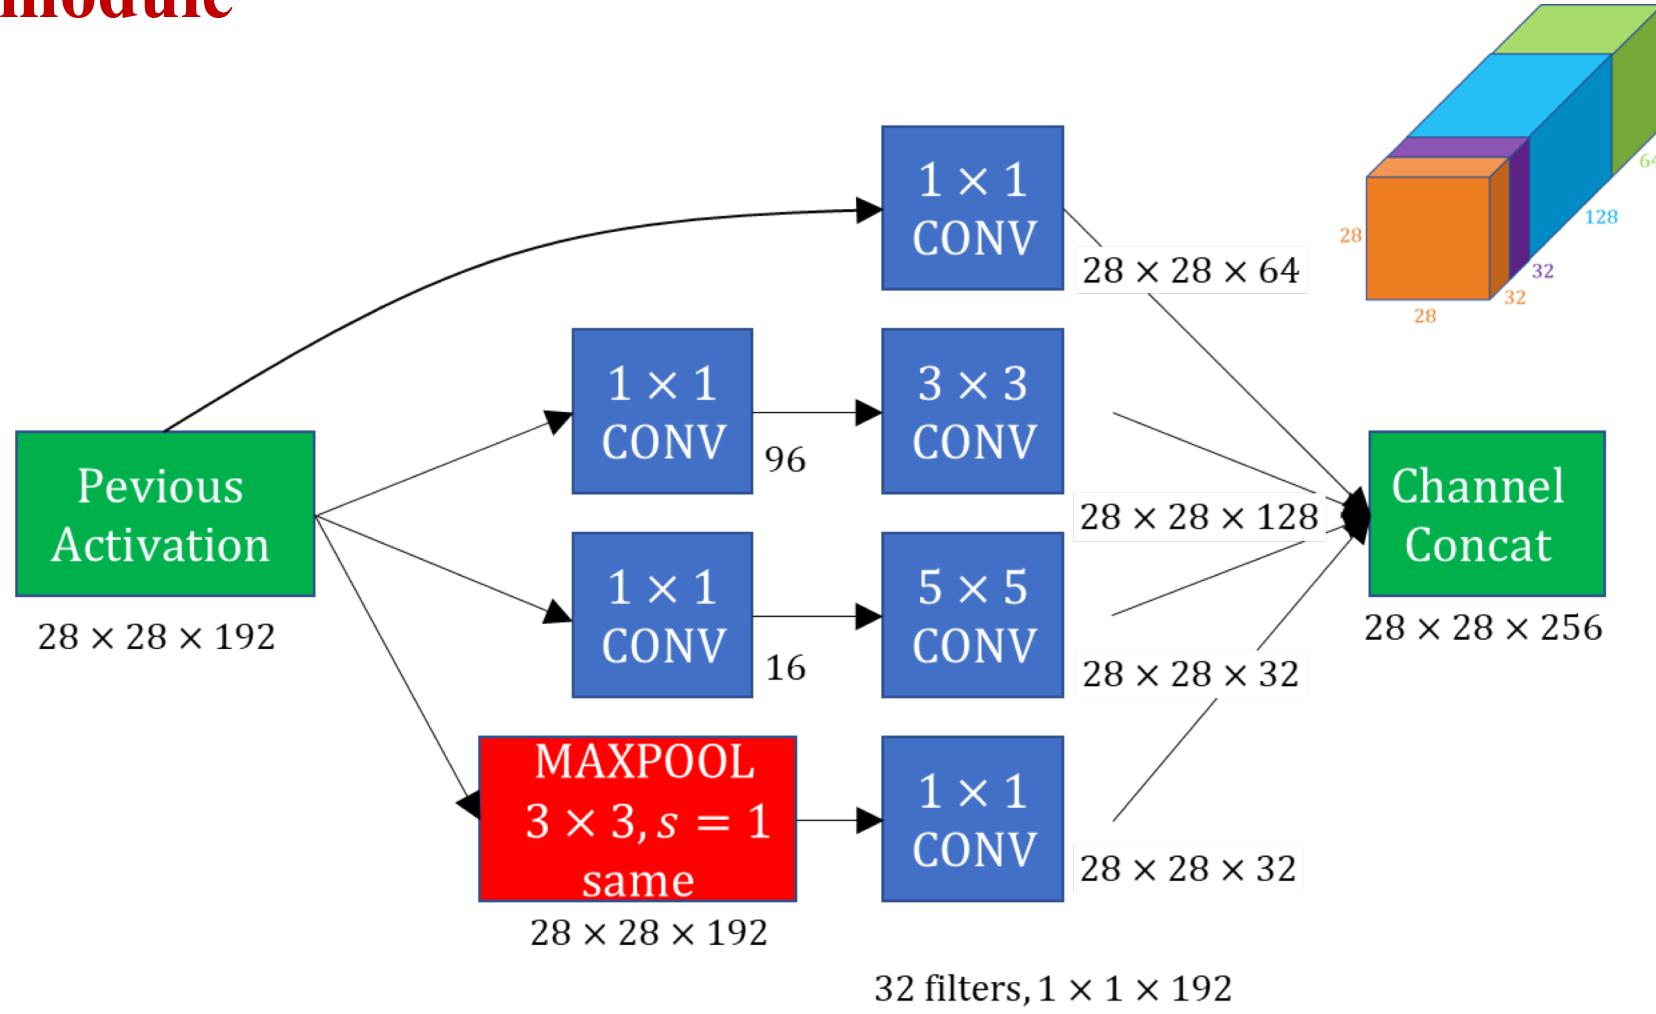
Do you want to use 1×1 filter, or 3×3 , or 5×5 ? Use a pooling layer? Or use them all?



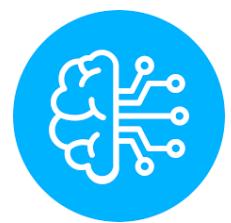
GoogleNet



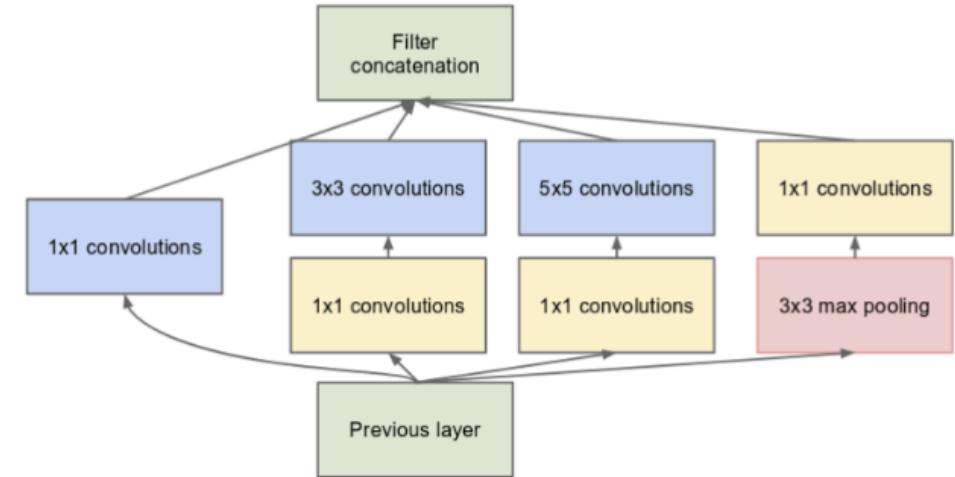
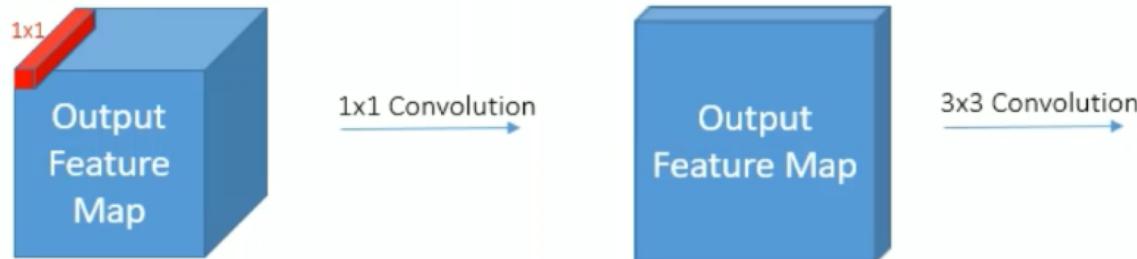
Inception module



GoogleNet



Reducing the depth while maintaining the size

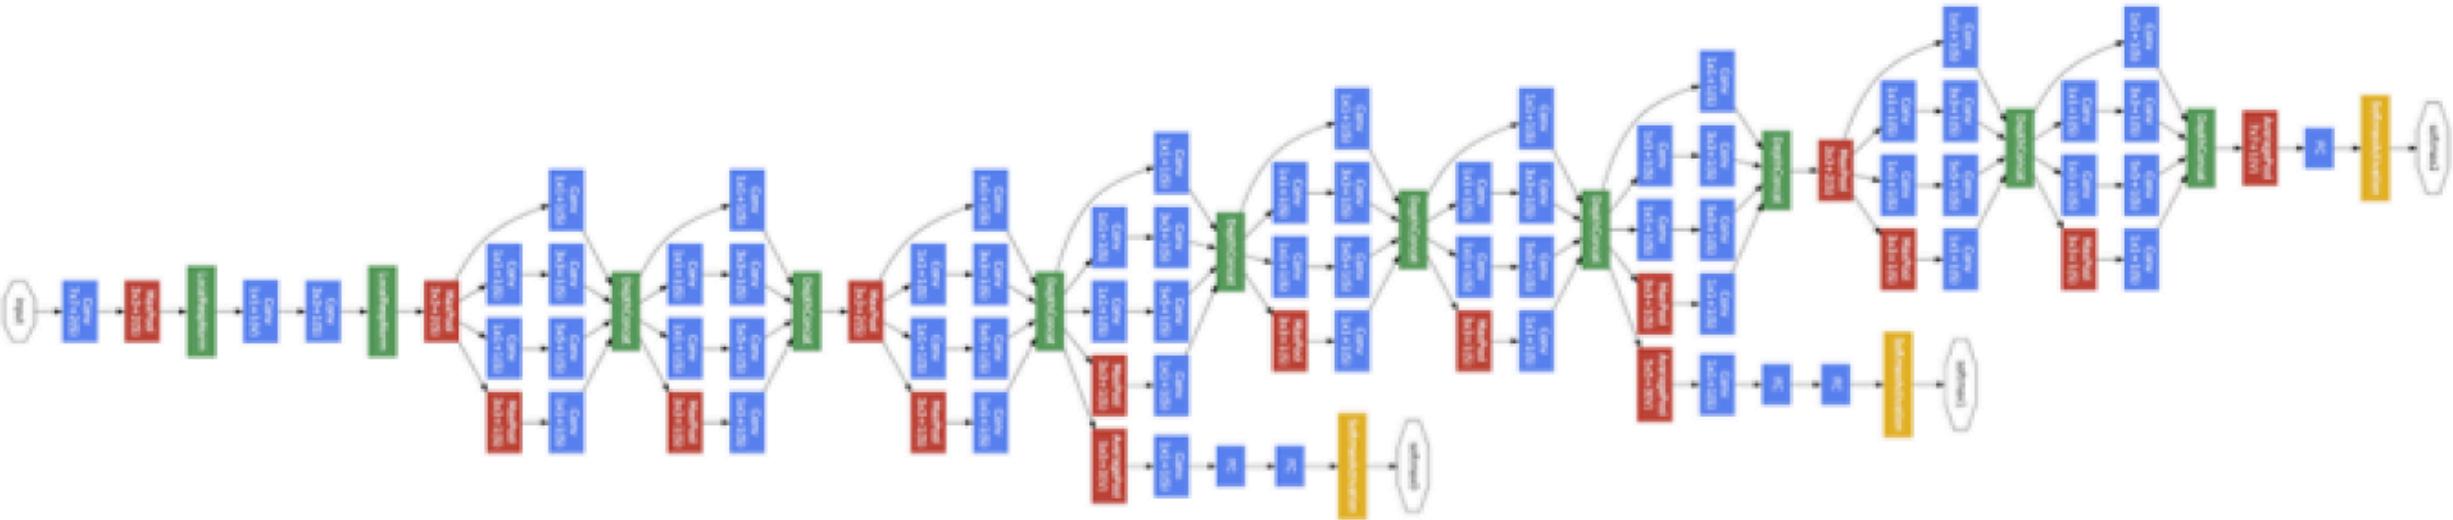
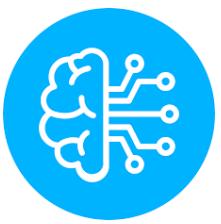


The 1x1 convolutions are a method for **dimensionality reduction**.

Example: Consider an input volume of **100x100x60**

Applying **30 filters** of 1x1 convolution would allow you to reduce the volume to **100x100x30** → The later convolution operations (such as 3x3 and 5x5 convolutions) won't have as large of a volume to deal with.

GoogleNet



The GoogleNet Architecture is 22 layers deep, with 27 pooling layers included. There are **9 inception modules** stacked linearly in total. The ends of the inception modules are connected to **the global average pooling layer**

Convolution
Pooling
Softmax
Other